



MICROECONOMETRICS

Bayesian Classifier and Decision Theory

WISE, XMU

Haihua Xie

27720181153991

Statistics

Oct. 20, 2019

1 Introduction

In this project, I'm going to introduce Bayesian classifier and decision theory. The bayesian decision theory will be introduced first, then the bayesian classifier, there are a large number of bayesian classifiers, we only cover the simple one: naive beysian classifier.

2 Bayesian decision theory

2.1 Bayes estimator for common loss functions

To make a decision, we have to choose an action a from some action space \mathcal{A} . Finally we incur some loss, $L(y, a)$, which measures how compatible our action a is with nature's hidden state y . Our goal is to devise a decision procedure or policy, $\delta : \mathcal{X} \rightarrow \mathcal{A}$, which specifies the optimal action for each possible input. By optimal, we mean the action that minimizes the expected loss: $\delta(\mathbf{x}) = \underset{a \in \mathcal{A}}{\operatorname{argmin}} \mathbb{E}[L(y, a)]$.

In the Bayesian approach to decision theory, the optimal action, having observed \mathbf{x} , is defined as the action a that minimizes the posterior expected loss : $\delta(\mathbf{x}) = \arg \min_{a \in \mathcal{A}} \rho(a|\mathbf{x})$.

The posterior expected loss: $\rho(a|\mathbf{x}) = \mathbb{E}_{p(y|\mathbf{x})}[L(y, a)] = \sum_y L(y, a)p(y|\mathbf{x})$. In the Bayesian version, we mean the expected value of y given the data we have seen so far. In the frequentist version, we mean the expected value of y and x that we expect to see in the future.

- **MAP estimate minimizes 0-1 loss**

The 0-1 loss is defined by:

$$L(y, a) = \mathbb{I}(y \neq a) = \begin{cases} 0 & \text{if } a = y \\ 1 & \text{if } a \neq y \end{cases}$$

The posterior expected loss is: $\rho(a|\mathbf{x}) = p(a \neq y|\mathbf{x}) = 1 - p(y|\mathbf{x})$. Hence the action that minimizes the expected loss is the MAP estimate: $y^*(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} p(y|\mathbf{x})$.

- **Posterior mean estimate minimizes ℓ_2 (quadratic)loss**

For continuous parameters, a more appropriate loss function is squared error loss, ℓ_2 loss ,or quadratic loss, defined as: $L(y, a) = (y - a)^2$. The posterior expected loss is given

by:

$$\rho(a|\mathbf{x}) = \mathbb{E}[(y - a)^2|\mathbf{x}] = \mathbb{E}[y^2|\mathbf{x}] - 2a\mathbb{E}[y|\mathbf{x}] + a^2$$

$$\hat{y} = \mathbb{E}[y|\mathbf{x}] = \int yp(y|\mathbf{x})dy$$

- **Posterior median minimizes ℓ_1 (absolute) loss**

The ℓ_1 (absolute) loss is defined by $L(y, a) = |y - a|$. The optimal estimate is the posterior median, i.e., a value of a such that $P(y < a|\mathbf{x}) = P(y \geq a|\mathbf{x}) = 0.5$.




Figure1: Graph of $L(y, a) = |y - a|^q$, for $q = 0.2$, $q = 1$ and $q = 2$

The ℓ_2 loss penalizes deviations from the truth quadratically, and thus is sensitive to outliers. A more robust alternative is the absolute or ℓ_1 loss.

- **Supervised learning**

Consider a prediction function $\delta : \mathcal{X} \rightarrow \mathcal{Y}$, and suppose we have some cost function $\ell(y, y')$, which gives the cost of predicting y' when the truth is y . The loss function is defined by

$$L(\boldsymbol{\theta}, \delta) = \mathbb{E}_{(\mathbf{x}, y) \sim p(\mathbf{x}, y|\boldsymbol{\theta})} [\ell(y, \delta(\mathbf{x}))] = \sum_{\mathbf{x}} \sum_y L(y, \delta(\mathbf{x})) p(\mathbf{x}, y|\boldsymbol{\theta})$$

This is known as the generalization error. Our goal is to minimize the posterior expected loss, given by:

$$\rho(\delta|\mathcal{D}) = \int p(\boldsymbol{\theta}|\mathcal{D}) L(\boldsymbol{\theta}, \delta) d\boldsymbol{\theta}$$

2.2 The false positive vs false negative tradeoff

Binary decision problem, such as hypothesis testing, two-class classification, object/event detection, etc.

The confusion matrix:

Table 1: Confusion matrix

	$y = 1$	$y = 0$	Sum
$\hat{y} = 1$	TP	FP	\hat{N}_+
$\hat{y} = 0$	FN	TN	\hat{N}_-
Sum	N_+	N_-	N

A general loss matrix

Table 2: Loss matrix

	$y = 1$	$y = 0$
$\hat{y} = 1$	0	L_{FP}
$\hat{y} = 0$	L_{FN}	0

where L_{FN} is the cost of a false negative, and L_{FP} is the cost of a false positive. The posterior expected losses for the two possible actions are given by

$$\rho(\hat{y} = 0|\mathbf{x}) = L_{FN} p(y = 1|\mathbf{x}) \text{ and } \rho(\hat{y} = 1|\mathbf{x}) = L_{FP} p(y = 0|\mathbf{x})$$

- ROC curves and Precision-recall curves

Accuracy: $\frac{TF+TN}{N}$;

Precision: $\frac{TP}{TP+FP}$;

Recall/sensitivity: $\frac{TP}{TP+FN} = \frac{TP}{N_+}$;

Specificity: $\frac{TN}{TN+FP} = \frac{TN}{N_-}$.

Table 3: Quantities derived from the confusion matrix

	$y = 1$	$y = 0$
$\hat{y} = 1$	TPR=sensitivity=recall	FPR=type I error rate
$\hat{y} = 0$	FNR=type II error rate	TNR=specificity

Rather than than computing the TPR and FPR for a fixed threshold τ , we can run our detector for a set of thresholds, and then plot the TPR vs FPR as an implicit function of τ . This is called a receiver operating characteristic or ROC curve. The quality of a ROC curve is often summarized as a single number using the area under the curve or AUC. Higher AUC scores are better; the maximum is obviously 1.




Figure2: ROC and Precision-Recall curves

(a) ROC curves for two hypothetical classification systems. A is better than B. We plot the true positive rate (TPR) vs the false positive rate (FPR) as we vary the threshold τ . We also indicate the equal error rate (EER) with the red and blue dots, and the area under the curve (AUC) for classifier B.

(b) A precision-recall curve for two hypothetical classification systems. A is better than B in this case.

• F-scores

For a fixed threshold, one can compute a single precision and recall value. These are often combined into a single statistic called the F score, or F1 score, which is the harmonic mean of precision and recall: estimate the precision and recall: $P = \frac{\sum_i y_i \hat{y}_i}{\sum_i \hat{y}_i}$, $R = \frac{\sum_i y_i \hat{y}_i}{\sum_i y_i}$; then F-score or F1-score:

$$F_1 = \frac{2}{1/P + 1/R} = \frac{2PR}{R + P} = \frac{2 \sum_{i=1}^N y_i \hat{y}_i}{\sum_{i=1}^N y_i + \sum_{i=1}^N \hat{y}_i}$$

To understand why we use the harmonic mean instead of the arithmetic mean, $(P + R)/2$, consider the following scenario: suppose we recall all entries, so $R = 1$, it's possible to get a very low precision, say, $p(y = 1) = 10^{-5}$, then the arithmetic mean is approximately 50%, while the harmonic mean of this strategy is approximately only 0.2%.

Macro-averaged F1, $\sum_{c=1}^C F_1(c)/C$ where $F_1(C)$ is the F1 score obtained on the task of distinguishing class c from all the others.

Micro-averaged F1, which is defined as the F1 score where we pool all the counts from each class' s contingency table.

Example:

Table 4: An example: macro-averaged vs. micro-averaged F1

Class 1			Class2			Pooled		
	$y = 1$	$y = 0$		$y = 1$	$y = 0$		$y = 1$	$y = 0$
$\hat{y} = 1$	10	10	$\hat{y} = 1$	90	10	$\hat{y} = 1$	100	20
$\hat{y} = 0$	10	970	$\hat{y} = 0$	10	890	$\hat{y} = 0$	20	1860

We see that the precision of class 1 is 0.5, and of class 2 is 0.9. The macro-averaged precision is therefore 0.7, whereas the micro-averaged precision is $100/(100 + 20) \approx 0.83$.

- **False discovery rates**

Suppose we are trying to discover a rare phenomenon using some kind of high throughput measurement device, such as a gene expression micro array, or a radio telescope. In such multiple hypotheses testing, we need to make many binary decision of the form $p(y_i = 1|\mathcal{D}) > \tau$, where $\mathcal{D} = \{\mathbf{x}_i\}_{i=1}^N$, we are classifying y_i based on all data, not just based on x_i , it's a simultaneous classification problem.

We are trying to minimize the expected number of the false positive to find an appropriate threshold τ .

$$FD(\tau, \mathcal{D}) = \sum_i (1 - p_i) \mathbb{I}(p_i > \tau)$$

$$FDR(\tau, \mathcal{D}) = FD(\tau, \mathcal{D})/N(\tau, \mathcal{D})$$

Where $p_i = p(y_i = 1|\mathcal{D})$ is your belief that this object exhibits the phenomenon in question. $N(\tau, \mathcal{D}) = \sum_i \mathbb{I}(p_i > \tau)$ is the number of discovered items. Given a desired FDR tolerance, say $\alpha = 0.05$, one can then adapt τ to achieve this, this is called the direct posterior probability approach to controlling the FDR. There are still a lot of researchers studying in the multiple hypotheses testing by controlling FDR, especially in the health and medicine.

3 Bayesian Classifier

Actually, the bayesian classifier is an application of the bayesian decision theory in solving some classification problems, I mainly introduce the naive bayes classifier, which is relatively simple and easy to understand. More complicated semi-naive bayes classifier and bayes network is left to be further reading.

3.1 Naive bayes classifier

Our goal is to classify vectors of discrete-valued features, $\mathbf{x} \in 1, 2, \dots, K^D$, where K is the number of values for each feature, and D is the number of features.

Assumption: the features are conditionally independent given the class label. This allows us to write the class conditional density as a product of one dimensional densities:

$$p(\mathbf{x}|y = c, \theta) = \prod_{j=1}^D p(x_j|y = c, \theta_{jc})$$

The resulting model is called a **naive bayes classifier**.

Gaussian naive bayes classifier. In the case of real-valued features, we can use the Gaussian distribution: $p(\mathbf{x}|y = c, \theta) = \prod_{j=1}^D \mathcal{N}(x_j|\mu_{jc}, \sigma_{jc}^2)$, where μ_{jc} is the mean of feature j in objects of class c , and σ_{jc}^2 is its variance.

Bernoulli naive bayes classifier. In the case of binary features, $x_j \in 0, 1$, we can use the Bernoulli distribution: $p(\mathbf{x}|y = c, \theta) = \prod_{j=1}^D \text{Ber}(x_j|\mu_{jc})$, where the μ_{jc} is the probability that feature j occurs in class c . This is sometimes called the multivariate Bernoulli naive Bayes model.

Categorical naive bayes classifier. In the case of categorical feature, $x_j \in 1, \dots, K$, we can use the multinoulli distribution: $p(\mathbf{x}|y = c, \theta) = \prod_{j=1}^D \text{Cat}(x_j|\mu_{jc})$, where the μ_{jc} is the histogram over the K possible values for x_j in class c .

3.2 Model fitting

To train a naive bayes classifier. This usually means computing the MLE or the MAP estimates of the parameters. The probability of single data case is given by:

$$p(\mathbf{x}_i, y_i|\boldsymbol{\theta}) = p(y_i|\boldsymbol{\pi}) \prod_j p(x_{ij}|\boldsymbol{\theta}_j) = \prod_c \pi_c^{\mathbb{I}(y_i=c)} \prod_j p(x_{ij}|\boldsymbol{\theta}_{jc})^{\mathbb{I}(y_i=c)}$$

Where π is the class prior. Hence, the likelihood is given by

$$\log p(\mathbf{x}_i, y_i|\boldsymbol{\theta}) = \sum_{c=1}^C N_c \log \pi_c + \sum_{j=1}^D \sum_{c=1}^C \sum_{i: y_i=c} \log p(x_{ij}|\boldsymbol{\theta}_{jc})$$

We see that this expression decomposes into a series of terms, one only concerns $\boldsymbol{\pi}$, and the other term contains $\boldsymbol{\theta}_{jc}$'s, then we can obtain the MLE for the class prior: $\hat{\pi}_c = \frac{N_c}{N}$, where the $N_c = \sum_i \mathbb{I}(y_i = c)$ is the number of examples in class c .

The MLE for the likelihood depends on the type of distribution we choose for each feature. For simplicity, suppose all features are binary, and $x_j|y = c \sim \text{Ber}(\theta_{jc})$, in this case, $\hat{\theta}_{jc} = \frac{N_{jc}}{N_c}$.