Frequentist Model Average Estimators

Nils Lid Hjort & Gerda Claeskens

Frequentist Model Average Estimators

Nils Lid HJORT and Gerda CLAESKENS

The traditional use of model selection methods in practice is to proceed as if the final selected model had been chosen in advance, without acknowledging the additional uncertainty introduced by model selection. This often means underreporting of variability and too optimistic confidence intervals. We build a general large-sample likelihood apparatus in which limiting distributions and risk properties of estimators post-selection as well as of model average estimators are precisely described, also explicitly taking modeling bias into account. This allows a drastic reduction in complexity, as competing model averaging schemes may be developed, discussed, and compared inside a statistical prototype experiment where only a few crucial quantities matter. In particular, we offer a frequentist view on Bayesian model averaging methods and give a link to generalized ridge estimators. Our work also leads to new model selection criteria. The methods are illustrated with real data applications.

KEY WORDS: Bias and variance balance; Growing models; Likelihood inference; Model average estimators; Model information criteria; Moderate misspecification.

1. INTRODUCTION AND SUMMARY

An impressive range of model selection criteria has been developed and refined over the past three decades. These have been constructed from different sets of intentions and have been aimed partly at general parametric models, whereas others have been geared toward special types of statistical models, such as time series, neural networks, and hazard rate regression; some are inspired by Bayesian considerations, whereas others are more traditional frequentistic; some have arisen via asymptotics and optimality properties for large samples, whereas others have been more fine tuned for moderate sample sizes; and so on. A fair number of these model choice schemes has also successfully made the passage from university blackboards to statistical software packages and the mainstream of applied statistical research. Methods such as the AIC and the BIC (the Akaike and the Bayesian information criteria), with suitable modifications, along with various stepwise methods for subset selection in regression models, are applied routinely also by nonspecialists. For overviews of model selection literature, one may consult the monograph by Burnham and Anderson (2002) and the introductory sections of Spiegelhalter, Best, Carlin, and van der Linde (2002) and Claeskens and Hjort (2003).

1.1 Estimator-Post-Selection Problems

It is fair to say, however, that far less work has been carried out, and even less has reached mainstream statistical applications, regarding the many complementary questions related to the consequences of model selection. In statistical practice one typically applies some off-the-shelf model selection scheme, perhaps supplemented with brief goodness-of-fit checking of residuals, to arrive at some "good model" that is thought to adequately reflect the main aspects of data—after which one proceeds with one's analysis as if this good model had been decided on in advance. It is clear that such analysis "hides (or ignores) some uncertainty"; reported confidence intervals tend to be too short, a hypothesis rejected at an announced 5% significance level might actually have been tested at a rather higher level, and so on. A central issue is that estimators formed after

model selection really are like mixtures of many potential estimators, namely, those that would have been computed had the random model selectors landed differently. A second theme is that it is sometimes advantageous to smooth estimators across several models, rather than sticking to only the model that is being reached by a single selection criterion.

There are at least two clear reasons fewer efforts have been devoted to these questions than to the primary ones related to finding "one good model." The first is that the selection strategies actually used by statisticians are difficult to describe accurately, as they involve many, partly nonformalized ingredients such as "looking at residuals" and "trying a suitable transformation." The second is that these questions of estimator post-selection behavior simply are harder to formalize and analyze.

An honorable exception is that of "Bayesian model averaging" (BMA), where more than a hundred papers have been published over the past decade. If a Bayesian can put down prior probabilities for a list of potential models, along with priors for the parameters of each model, then the Bayesian machinery is, in principle, capable of delivering the posterior distribution of any interest parameter (provided it retains a precise interpretation and is well defined across the models under study). The tutorial by Hoeting, Madigan, Raftery, and Volinsky (1999) discusses pertinent issues of interpretation and implementation via the machinery of Markov chain Monte Carlo (MCMC), where the chains in question move between models of different dimensions; see also Green (2003) for a review of transdimensional MCMC theory. With BMA methodology the extra estimator variability stemming from not knowing the correct model a priori is adequately taken into account.

The approach remains problematic, however. First, there are difficulties associated with the often ad hoc way in which the prior probabilities for a (sometimes long) list of models is set up; see the discussion to Hoeting et al. (1999). Second, we raise concern for the fact that the typical application of BMA involves mixing together many conflicting prior opinions regarding interest parameters. If μ is some parameter of interest, and $\mu = \mu(\alpha_j)$ in terms of the parameters α_j of candidate model j, with prior $\pi_j(\alpha_j)$, this leads to a prior $\bar{\pi}_j(\mu)$, say; why would a statistician entertain many different such priors inside the same problem formulation, and what are the consequences in cases where some of these have clear clashes? Finally, even

though BMA "works," insofar as adequate analysis of data can be carried out after judicious selection of models, prior probabilities for these, and prior densities for parameters in each model, rather little appears to be known about the actual performance or behavior of the consequent inferences, such as estimator precision.

The present article aims at establishing a framework where properties of estimator-post-selection and estimator average methods can be accurately described. Our framework is general and unified and involves large-sample likelihood approximations across a list of parametric models. The end result is a machinery for "frequentist model averaging" (FMA), to be partly contrasted with that of BMA. Within this context many natural model averaging strategies can be developed and compared. Our results also shed light on the behavior of BMA schemes, in fact, by leading to precise large-sample results about their behavior.

1.2 An Illustration: Averages Over Logistic Regressions

To illustrate and pinpoint some of the problems associated with model selection and model averaging, consider the following example. The dataset studied is taken from appendix I in Hosmer and Lemeshow (1989) and concerns factors that may influence the birth weight of babies, in particular, the event that the baby weighs less than 2,500 grams. Covariate information for the $n = 189$ mothers in question included weight just prior to pregnancy (x_2, in pounds), age (x_3), as well as indicators for race "black" (x_4) and race "other" (x_5); mothers with $x_4 = 0$ and $x_5 = 0$ are of race "white." For the purposes of this article, we make the assumption that

$$p(x, u) = \Pr\{\text{low birth weight} \mid x, u\} = \frac{\exp(x^t\beta + u^t\gamma)}{1 + \exp(x^t\beta + u^t\gamma)},$$

where $x = (1, x_2)^t$ is always to be included in the logistic regression, whereas subsets of $u = (x_3, x_4, x_5)^t$ may or may not enter the equation. See also Claeskens and Hjort (2003).

It is convenient to label the eight potential submodels "0", "3", "4", "5", "34", "35", "45", and "345", corresponding to inclusion or exclusion of these three extra covariates. We shall take an interest in estimating three parameters, the probability of low birth weight for the average "white" and "black" mothers and for the ratio of these two. Table 1 gives estimates along with associated standard errors for these three estimands for each of the eight possible models. The table also includes mi-

nus AIC and minus BIC, where AIC is twice the maximized log-likelihood minus say $2k$, where k is the number of parameters in the model, whereas the BIC is twice the maximized log-likelihood minus $k \log n$. We see that the AIC selects "4" ahead of "45," whereas the BIC prefers the narrow model "0" ahead of "4." See also Claeskens and Hjort (2003) for further analysis of these data using the focused information criterion (FIC), which finds the best model for a given interest parameter.

The estimated standard deviations given here have been computed via familiar delta method algebra and approximate normality of the maximum likelihood estimators, and under the typical assumption that the model under consideration is adequate. Although the sampling variance perhaps may be adequately estimated here (conditional on the model), there is potential modeling bias, not reflected in the table and not easy to assess. Our article will develop methods that, in particular, make it possible to answer the following questions:

1. If a statistician uses the estimators dictated by the AIC (here .269, .412, and 1.533 for the three parameters), what are the real variances of these, and what are the biases stemming from the modeling imperfections of the selected logistic equation? How trustworthy are the confidence intervals delivered by standard use?

2. Similarly, if another statistician uses the BIC to decide on estimators (here .298, .256, and .861), how large might the modeling biases be, and what are the real variances involved?

3. Are there advantages to taking suitable averages across models, for example, weighted averages over those with best AIC, BIC, or FIC scores? What, then are the biases and variances involved? How can adequate confidence intervals be constructed?

4. When will the simple "narrow method," which here corresponds to disregarding the extra covariates, be more accurate than the "full model method," which includes all five logistic parameters in the inference?

5. Could it be advantageous here to trust covariate x_2 fully (along with $x_1 = 1$), but to trust the influence of x_3, x_4, and x_5 less, in the sense of shrinking estimated logistic coefficients for these three toward 0?

6. If a BMA regime is used here, what is its (frequentist) behavior, and how do different BMA schemes compare in performance?

7. Are there FMA schemes with suitable optimality properties?

Table 1. For Submodels Corresponding to Inclusion or Not of the Covariates x_3, x_4, and x_5, the Table Lists Minus AIC, Minus BIC, and Then Estimates Along With Estimated Standard Deviations (computed under the model assumption in question). These are the low-birth-weight probabilities p(white), p(black), and the ratio p(black)/ p(white)

Model	−AIC	−BIC	White	SE	Black	SE	Ratio	SE
0	232.691	239.174[*]	.298	.035	.256	.040	.861	.060
3	233.123	242.849	.288	.035	.272	.043	.945	.094
4	231.075[*]	240.800	.269	.037	.412	.101	1.533	.423
5	234.101	243.826	.279	.041	.242	.043	.868	.062
34	232.175	249.068	.264	.037	.413	.101	1.564	.435
35	234.677	247.644	.272	.041	.259	.046	.950	.097
45	231.259	244.226	.231	.044	.414	.100	1.794	.547
345	232.661	248.869	.230	.044	.414	.100	1.801	.551

NOTE: The asterisk indicates the selected model.

1.3 Related Work

As mentioned previously, the Bayesian literature so far decidedly outgoliaths its frequentist counterpart concerning model averaging inference and estimator postselection performance. Some work in the frequentist directions has, however, been done over the last decade.

Hurvich and Tsai (1990) pointed out that for linear regression models coverage rates of confidence intervals for regression parameters, conditional on the selected model, are much smaller than the nominal coverage rates. Such problems have been further addressed by Chatfield (1995) and Draper (1995). Also, in a linear regression setting, in the presence of a finite-dimensional nuisance parameter, Kabaila (1995, 1998) considered the effect of model selection on the construction of confidence intervals as well as on prediction intervals. Pötscher (1991) considered a sequence of nested models containing an increasing number of parameters $\theta_1, \ldots, \theta_q$ and possibly a nuisance parameter η in which a backward model selection is performed. He made the assumption that there is a true model containing parameters $(\eta, \theta_1, \ldots, \theta_{q_0})$, where $1 \leq q_0 \leq q$. Leeb and Pötscher (2000) further built on this subject and obtained distributions of post-model selection estimators under the condition of possibly selecting an incorrect model with fewer than q_0 parameters. Their methods are restricted to linear regression models $Y = X\theta + \varepsilon$ with independent and identically distributed Gaussian error terms, and for a similar backward selection procedure, employing a t test at each stage of the selection procedure. Also, in linear models, Sen and Saleh (1987) studied the asymptotic distribution after a preliminary test for the presence of part of the regression coefficients, hence dealing with two possible models for the data. Bühlmann (1999) investigated the local consistency of post-model selection estimators under a set of conditions that imply all "local" models have the same dimension asymptotically.

A few non-Bayesian methods for model averaging have been proposed in the literature. There is, of course, a large literature on model selection methods, which can be considered as hard-threshold averages; see Claeskens and Hjort (2003). George (1986a,b) investigated multiple-shrinkage estimators in the normal model. Also, Foster and George (1994) explicitly analyzed performance of post-selection estimators in a normal regression context. Rao and Tibshirani (1997) constructed an out-of-bootstrap method that leaves out one training point and constructs bootstrap model weights depending on how well the remaining bootstrap data predict the left-out value. They did not provide any asymptotic distribution theory for the model averaging estimator. The adaptive regression by mixing of Yang (2001) splits the dataset into two parts, where one part is used for estimation and the other for measuring the quality of predictions, on the basis of which model weights are constructed. Buckland, Burnham, and Augustin (1997) constructed model averaging weights based on the values of the AIC or BIC scores, further discussed in Burnham and Anderson (2002, chap. 6). This may accordingly be seen as suggestions for problem 3 described at the end of Section 1.2. The construction of Buckland et al. is somewhat ad hoc, however, and they do not really analyze the performance of the resulting estimator. The results of Section 4 can be used to accurately describe its behavior, and also answer other questions raised in their article.

1.4 The Present Article

In Section 2 we build and discuss a general model selection framework, involving a finite number of parametric extensions around a given parametric basis model. This includes problems of subset selection in general regression models and, with further amendments, situations with memory order and averaging order for time series and Markov chain models. Section 3 develops the necessary theory of maximum likelihood estimators inside such a framework, where modeling bias is explicitly present and taken into account for each candidate estimator. Some attention is given to the behavior of the AIC criterion. In Section 4 we describe a fairly general class of model average estimators, which compromise across a set of candidate models, and derive their limit distributions. These are not normal, but rather nonlinear mixtures of normals. This is, in particular, true for estimator-after-selection schemes. We also pinpoint how the confidence level of the confidence intervals becomes lower than the "intended level" when the model selection step is being ignored. Various natural FMA strategies are proposed in Section 5, including Bayesian and empirical Bayesian variants. Then we illustrate our FMA machinery for some applications in Section 6. Section 7 extends the class of compromise estimators further, allowing "generalized ridging," where estimates of potential extensions of a given model are being shrunk in a controlled fashion. This may often lead to smaller variances without significantly increased sizes of modeling bias. Then we turn to a frequentist view of the BMA methods in Section 8. Apparently, despite a flurry of BMA activity over the last decade, the performance of BMA schemes has not been studied in the classical sense of limiting distributions and large-sample approximations to risks; we do so here. In Section 9 we give some brief analysis of risk behavior and comparisons, applying and illustrating the theoretical results of previous sections. Our article ends with a list of concluding remarks in Section 10, some pointing to further research. All proofs are given in the Appendix.

2. A MODEL AVERAGING FRAMEWORK

This section establishes a fruitful general framework for model choice and model average estimators. The motivation is to start with a "narrow" model, perhaps of standard type, and then add one or more additional parameters to be able to reflect further features of the data generating mechanisms at work. This section partly parallels Section 2 of Claeskens and Hjort (2003), which focuses on model selection, whereas we are also concerned here with model averaging.

2.1 Models With iid Data

Suppose independent data $Y_1, \ldots, Y_n$ come from density f. Inference is sought for a certain parameter estimand $\mu = \mu(f)$. We start with a narrow model of the type $f(y, \theta)$ with a p vector of parameters θ. The extended models take the form $f(y, \theta, \gamma)$ with an additional q vector of γ parameters, where $\gamma = \gamma_0$ corresponds to the narrow model in the sense that $f(y, \theta) = f(y, \theta, \gamma_0)$. Thus, γ_0 is fixed and known. Here one may consider employing suitable submodels, corresponding to having some of the γ_j parameters equal to $\gamma_{j,0}$ whereas others are not. Using a bigger model would typically mean less

modeling bias but increased estimation variance and vice versa. At the outset there are 2^q such submodels to consider, one for each subset S of $\{1, \ldots, q\}$.

In this framework there are a variety of estimators to consider, ranging from $\hat{\mu}_{\text{full}} = \mu(\hat{\theta}_{\text{full}}, \hat{\gamma}_{\text{full}})$ using maximum likelihood estimators in the widest model where $S = \{1, \ldots, q\}$ to the simpler $\hat{\mu}_{\text{narr}} = \mu(\hat{\theta}_{\text{narr}}, \gamma_0)$, which employs maximum likelihood estimation in the narrow model where $S = \varnothing$. The general submodel estimator is

$$\hat{\mu}_S = \mu(\hat{\theta}_S, \hat{\gamma}_S, \gamma_{0, S^c}), \qquad \text{where} \quad S \subset \{1, \ldots, q\}, \quad (2.1)$$

found via maximum likelihood in the model that includes exactly the γ_j parameters for $j \in S$ while keeping the others at $\gamma_{0,j}$ (S^c is the complement of S). The narrow model corresponds to S being the empty set. Further special cases would be the nested ones corresponding to $S = \{1, \ldots, k\}$ for $k = 1, \ldots, q$.

Our intention is to investigate what happens to all the $\hat{\mu}_S$ estimators, and, importantly, averaged versions of these, in the local misspecification framework

$$f_{\text{true}}(y) = f_n(y) = f(y, \theta_0, \gamma_0 + \delta/\sqrt{n}). \quad (2.2)$$

The $\delta_1, \ldots, \delta_q$ parameters signify the degrees of the model departures in directions $1, \ldots, q$, with due influence on the estimand $\mu_{\text{true}} = \mu(\theta_0, \gamma_0 + \delta/\sqrt{n})$. Later on we give results for the limiting risk of estimators $\sum_S c(S) \hat{\mu}_S$, with random weights summing to 1. We assume that $\mu(\theta, \gamma)$ is smooth in a neighborhood of (θ_0, γ_0). The aim is to understand and assess large-sample approximations to distributions and say mean squared errors of subset and model average estimators in situations where the data come from $f(y, \theta, \gamma)$ with γ not too far from γ_0, and it is for this reason that we work under the (2.2) scenario. The $O(1/\sqrt{n})$ framework chosen here is canonical in the sense that it leads to the most fruitful large-sample approximations, with squared model biases and estimator variances as exchangeable currencies, both of size $O(1/n)$.

Our framework amounts to studying perturbations around a given narrow model in certain directions, expressed in (2.2) by letting γ vary around γ_0, and various consequences for different estimators are highlighted and discussed in the following sections. One may wonder whether further consequences of importance would emerge if we, in addition, perturb the θ part of the model around the null value θ_0. It turns out that there is no additional gain in considering such scenarios, as judged by what is to be learned from large-sample approximations of estimators and their performances; see Remark 4.1.

2.2 Subset Selection and Mixtures of Regression Models

With some efforts the preceding framework may be generalized to encompass regression situations. Suppose Y_i for given covariate vectors $x_i = (x_{i,1}, \ldots, x_{i,p})^t$ are independent, with density of the type

$$f_{i,\text{true}}(y \mid x_i) = f(y \mid x_i, \beta_0, \sigma_0, \gamma_0 + \delta/\sqrt{n}),$$

most often with a p-dimensional β parameter, a scale parameter σ, and up to q further parameters γ. These "extra" parameters could be associated with interactions among the $x_{i,j}$ covariates or with other regressors. They could also help describe aspects of the variance structure, such as a parametric model for the conditional variance in linear regression. It is

again required that $\gamma = \gamma_0$ leads back to the narrow model with only β and σ present. Focus parameters of interest take the form $\mu = \mu(\beta, \sigma, \gamma)$, which here corresponds to $\mu_{\text{true}} = \mu(\beta_0, \sigma_0, \gamma_0 + \delta/\sqrt{n})$. This could be the median regression surface or the standard deviation function evaluated at a point x_0, a quotient between two regression coefficients or between two values of the mean regression function, and so on. In this framework we may consider submodel estimators $(\hat{\beta}_S, \hat{\sigma}_S, \hat{\gamma}_S)$ via maximum likelihood in the model that employs γ_j's for $j \in S$. This leads to the estimator $\hat{\mu}_S = \mu(\hat{\beta}_S, \hat{\sigma}_S, \hat{\gamma}_S)$ for the focus parameter. In Section 4 we give results for limiting risks of model average estimators $\sum_S c(S) \hat{\mu}_S$.

The type of local neighborhood models described here also have parallels in time series and Markov chains, where it could be advantageous to weight across models with different memory lengths.

3. LIMIT DISTRIBUTION THEORY

In this section we establish the notation necessary for handling analysis in the various submodels and then sort out the behavior of different maximum likelihood estimators. We also give relevant limit results for log-likelihood ratios, which, in particular, are needed to understand the performance characteristics of the AIC model choice criterion.

We work throughout under traditional conditions of regularity, sufficient to apply familiar likelihood asymptotics arguments, as laid out, for example, in Lehmann (1983, chap. 6). Thus, the log density admits two continuous partial derivatives in all directions; (θ_0, γ_0) is an inner point of the parameter space; the variance matrix of the score function statistic is finite and positive definite in a neighborhood around this null point; and certain derivative operations can be taken under the integral sign. Details and proofs of the lemmas are given in the Appendix.

3.1 Notation for Calculus in Submodels

Consider the score function

$$\begin{pmatrix} U(y) \\ V(y) \end{pmatrix} = \begin{pmatrix} \partial \log f(y, \theta_0, \gamma_0)/\partial \theta \\ \partial \log f(y, \theta_0, \gamma_0)/\partial \gamma \end{pmatrix},$$

with a p-dimensional U and q-dimensional V. Their $(p+q) \times (p+q)$ variance matrix at the null model is

$$J_{\text{full}} = \begin{pmatrix} J_{00} & J_{01} \\ J_{10} & J_{11} \end{pmatrix}, \quad \text{with inverse} \quad J_{\text{full}}^{-1} = \begin{pmatrix} J^{00} & J^{01} \\ J^{10} & J^{11} \end{pmatrix},$$

say; in particular, let

$$K = J^{11} = (J_{11} - J_{10} J_{00}^{-1} J_{01})^{-1}.$$

Under consideration are models indexed by subsets S of $\{1, \ldots, q\}$. We let π_S be the projection matrix mapping $v = (v_1, \ldots, v_q)^t$ to the subvector $\pi_S v = v_S$ of components v_j with $j \in S$. Hence, π_S is of size $|S| \times q$ with $|S|$ being the size of S. For $V_S = \pi_S V$ we then have

$$J_S = \text{Var}_0 \begin{pmatrix} U(Y) \\ V_S(Y) \end{pmatrix} = \begin{pmatrix} J_{00} & J_{01,S} \\ J_{10,S} & J_{11,S} \end{pmatrix}$$

$$= \begin{pmatrix} J_{00} & J_{01} \pi_S^t \\ \pi_S J_{10} & \pi_S J_{11} \pi_S^t \end{pmatrix}.$$

We shall also need its inverse matrix, which has blocks $J^{11,S} = (\pi_S K^{-1} \pi_S^{\mathrm{t}})^{-1} = K_S$, $J^{01,S} = -J_{00}^{-1} J_{01} \pi_S^{\mathrm{t}} K_S$, and $J^{00,S} = J_{00}^{-1} + J_{00}^{-1} J_{01} \pi_S^{\mathrm{t}} K_S \pi_S J_{10} J_{00}^{-1}$.

Lemma 3.1. Consider the averages $\overline{U}_n = n^{-1} \sum_{i=1}^{n} U(Y_i)$ and $\overline{V}_n = n^{-1} \sum_{i=1}^{n} V(Y_i)$. Under the sequence of local alternatives (2.2),

$$\begin{pmatrix} \sqrt{n}\,\overline{U}_n \\ \sqrt{n}\,\overline{V}_n \end{pmatrix} \xrightarrow{d} \begin{pmatrix} J_{01}\delta \\ J_{11}\delta \end{pmatrix} + \begin{pmatrix} M \\ N \end{pmatrix},$$

where $\quad \begin{pmatrix} M \\ N \end{pmatrix} \sim N_{p+q}(0, J_{\mathrm{full}})$.

3.2 Behavior of Maximum Likelihood Estimators in Submodels

Let $(\hat{\theta}_S, \hat{\gamma}_S)$ denote maximum likelihood estimators in the model that includes γ_j parameters for $j \in S$.

Lemma 3.2. Under the sequence of models f_{true} of (2.2),

$$\begin{pmatrix} \sqrt{n}(\hat{\theta}_S - \theta_0) \\ \sqrt{n}(\hat{\gamma}_S - \gamma_{0,S}) \end{pmatrix} \xrightarrow{d} \begin{pmatrix} C_S \\ D_S \end{pmatrix} = J_S^{-1} \begin{pmatrix} J_{01}\delta + M \\ \pi_S J_{11}\delta + N_S \end{pmatrix}$$

$$\sim N_{p+|S|} \left(J_S^{-1} \begin{pmatrix} J_{01} \\ \pi_S J_{11} \end{pmatrix} \delta, J_S^{-1} \right).$$

Before stating the next lemma, it is convenient to introduce some more notation, which also will be needed later. Define first $W = J^{10}M + J^{11}N = K(N - J_{10}J_{00}^{-1}M)$. Here $M \sim N_p(0, J_{00})$, and it is not difficult to establish that M and W are stochastically independent, with W having a $N_q(0, K)$ distribution. It follows from Lemma 3.2 and a little algebra that $\hat{\delta}_S = \sqrt{n}(\hat{\gamma}_S - \gamma_{0,S})$ tends in distribution to $D_S = K_S \pi_S K^{-1}(\delta + W)$. In particular,

$$D_n = \hat{\delta}_{\mathrm{full}} = \sqrt{n}(\hat{\gamma}_{\mathrm{full}} - \gamma_0) \xrightarrow{d} D = \delta + W \sim N_q(\delta, K). \quad (3.1)$$

Next let

$$H_S = K^{-1/2} \pi_S^{\mathrm{t}} K_S \pi_S K^{-1/2},$$
$$\omega = J_{10} J_{00}^{-1} \frac{\partial \mu}{\partial \theta} - \frac{\partial \mu}{\partial \gamma}, \quad (3.2)$$

where the partial derivatives indicated are evaluated at the null model (θ_0, γ_0). Note that ω is determined by the specifics of the focus parameter μ. The H_S is a $q \times q$ projection matrix, being symmetric and idempotent, and is orthogonal to $I - H_S$. We define $H_\varnothing$ as the null matrix of size $q \times q$.

Lemma 3.3. Assume $\mu(\theta, \gamma)$ has continuous partial derivatives in a neighborhood of (θ_0, γ_0). Then the maximum likelihood estimator of μ in the S model has limiting distribution of the form

$$\sqrt{n}(\hat{\mu}_S - \mu_{\mathrm{true}})$$

$$\xrightarrow{d} \Lambda_S = \left(\frac{\partial \mu}{\partial \theta} \right)^{\mathrm{t}} J_{00}^{-1} M + \omega^{\mathrm{t}}(\delta - K^{1/2} H_S K^{-1/2} D),$$

where the partial derivatives indicated are evaluated at the null model (θ_0, γ_0). The limiting variable is normal, with mean $\omega^{\mathrm{t}}(I - K^{1/2} H_S K^{-1/2})\delta$ and variance $(\frac{\partial \mu}{\partial \theta})^{\mathrm{t}} J_{00}^{-1} \frac{\partial \mu}{\partial \theta} + \omega^{\mathrm{t}} K^{1/2} H_S K^{1/2} \omega$.

3.3 AIC Calculus

The Akaike information criterion is equal to

$$\mathrm{AIC}_{n,S} = 2\sum_{i=1}^{n} \log f(Y_i, \hat{\theta}_S, \hat{\gamma}_S, \gamma_{0,S^c}) - 2|S|,$$

again with $|S|$ being the number of elements in S. Its typical use is to pick out the model with the largest value of this criterion. To understand the behavior of this criterion in the present framework, we start out with the likelihood ratio statistic, expanding it to the second order, using familiar arguments. This leads to

$$G_{n,S} = 2\sum_{i=1}^{n} \log\{ f(Y_i, \hat{\theta}_S, \hat{\gamma}_S, \gamma_{0,S^c})/f(Y_i, \theta_0, \gamma_0) \}$$

$$\doteq n \begin{pmatrix} \overline{U}_n \\ \overline{V}_{n,S} \end{pmatrix}^{\mathrm{t}} J_S^{-1} \begin{pmatrix} \overline{U}_n \\ \overline{V}_{n,S} \end{pmatrix}$$

$$\xrightarrow{d} \begin{pmatrix} J_{01}\delta + M \\ \pi_S J_{11}\delta + N_S \end{pmatrix}^{\mathrm{t}} J_S^{-1} \begin{pmatrix} J_{01}\delta + M \\ \pi_S J_{11}\delta + N_S \end{pmatrix}.$$

(Here and later we use for notational simplicity $X_n \doteq X_n'$ to indicate that the difference between the two variables tends to 0 in probability; thus they have the same limit distribution, if it exists.) This is a noncentral chi-squared with $p + |S|$ degrees of freedom. Furthermore,

$$G_{n,S} - G_{n,\varnothing} = n(\overline{V}_{n,S} - J_{10,S} J_{00}^{-1} \overline{U}_n)^{\mathrm{t}}$$
$$\times J^{11,S}(\overline{V}_{n,S} - J_{10,S} J_{00}^{-1} \overline{U}_n)$$
$$\xrightarrow{d} (K_S^{-1}\delta + N_S - J_{10,S} J_{00}^{-1} M)^{\mathrm{t}}$$
$$\times K_S(K_S^{-1}\delta + N_S - J_{10,S} J_{00}^{-1} M),$$

which is a noncentral $\chi^2_{|S|}(\delta^{\mathrm{t}} K_S^{-1} \delta)$.

Using a combination of previous arguments, $\hat{\delta}_S = \sqrt{n}(\hat{\gamma}_S - \gamma_{0,S})$ is at most $o_p(1)$ away from $\sqrt{n} K_S \pi_S (\overline{V}_n - J_{10} J_{00}^{-1} \overline{U}_n)$, and, similarly, $D_n \doteq \sqrt{n} K(\overline{V}_n - J_{10} J_{00}^{-1} \overline{U}_n)$, which implies $\hat{\delta}_S \doteq K_S \pi_S K^{-1} D_n$. The important consequence is that $\hat{\gamma}_S$, the estimator based on the S subset model, can be expressed, within the first-order local asymptotic framework, as a function of $\hat{\gamma}_{\mathrm{full}}$. It also follows that the AIC criterion can be expressed in terms of D_n as

$$\mathrm{AIC}_{n,S} = G_{n,S} - G_{n,\varnothing} - 2|S|$$
$$= D_n^{\mathrm{t}} K^{-1/2} H_S K^{-1/2} D_n - 2|S| + o_p(1). \quad (3.3)$$

3.4 Results for the Regression Framework

The methods and results given previously generalize to the regression-type framework of Section 2.2 without too many difficulties, with the appropriate modifications and regularity conditions. An important ingredient is

$$J_{n,\mathrm{full}} = \frac{1}{n}\sum_{i=1}^{n} \mathrm{Var}_0 \begin{pmatrix} \partial \log f(Y_i \mid x_i, \beta_0, \sigma_0, \gamma_0)/\partial \beta \\ \partial \log f(Y_i \mid x_i, \beta_0, \sigma_0, \gamma_0)/\partial \sigma \\ \partial \log f(Y_i \mid x_i, \beta_0, \sigma_0, \gamma_0)/\partial \gamma \end{pmatrix}$$

$$= \begin{pmatrix} J_{n,00} & J_{n,01} \\ J_{n,10} & J_{n,11} \end{pmatrix},$$

say, where $J_{n,00}$ is of size $(p+1) \times (p+1)$ and $J_{n,11}$ of size $q \times q$. This matrix is assumed to converge to a suitable J_{full} as n increases. There are natural analogs of Lemmas 3.1–3.3 as well as for the AIC calculus results. Concrete regularity conditions would depend on the regression models studied. They would typically include assumptions of the Lindeberg–Lyapunov type $n^{-1/2} \max_{i \le n} \|x_i\| \to 0$, which are fulfilled in situations where the x_i's come from some covariate distribution with finite second moment.

4. ESTIMATORS AFTER SELECTION AND COMPROMISE ESTIMATORS

The estimator employed by a statistician using a model selection criterion really takes the form $\hat{\mu} = \hat{\mu}_{\widehat{S}}$, where $\widehat{S}$ is the (random) set picked out by the selection procedure, for example, the one exhibiting the largest $\text{AIC}_{n,S}$ number. The behavior of a large class of such mixed-situation estimators, which we may think of as frequentist model average estimators, is studied in this section. Our results are, in particular, used to pinpoint the overoptimistic nature of traditionally employed confidence intervals wrt coverage probability.

4.1 Compromise Estimators

To be able to single out submodels with more influence than others, it is natural to employ $\hat{\delta}_S = \sqrt{n}(\hat{\gamma}_S - \gamma_{0,S})$ in a suitable form. We saw in Section 3.3 that the behavior of $\hat{\delta}_S$ is essentially determined by that of $D_n = \hat{\delta}_{\text{full}}$ of (3.1). Which submodel is picked out by the AIC method, for example, is determined by D_n; see (3.3). This motivates studying the large class of compromise estimators, those taking the form

$$\hat{\mu} = \sum_S c(S \mid D_n)\hat{\mu}_S, \qquad (4.1)$$

where the sum is potentially over all subsets of $\{1, \ldots, q\}$, including the empty subset, which corresponds to the narrow model. The weight functions $c(S \mid d)$ are required to sum to 1 for each d, because otherwise the estimator is not consistent. A special case would be $\hat{\mu} = \sum_{k=0}^{q} c(\{1, \ldots, k\} \mid D_n)\hat{\mu}_k$, say, indicating a mixture of estimators $\hat{\mu}_k$ constructed from the model with $S = \{1, \ldots, k\}$. Estimators formed after using the AIC criterion for nested submodels would be of this type, for example.

For a general compromise estimator of type (4.1), and with H_S as in (3.2), define $G(d) = K^{-1/2}\{\sum_S c(S \mid d)H_S\}K^{1/2}$ and

$$\hat{\delta}(D) = G(D)^t D = K^{1/2}\left\{\sum_S c(S \mid D)H_S\right\}K^{-1/2}D. \quad (4.2)$$

Then $G(d)$ is a $q \times q$ matrix of functions in $d = (d_1, \ldots, d_q)^t$, and $\hat{\delta}(D)$ is to be seen as an estimator of δ based on D. Recall that $D_n = \hat{\delta}_{\text{full}}$ tends to $D \sim N_q(\delta, K)$ by (3.1).

Theorem 4.1. As long as the weight functions $c(S \mid d)$ sum to 1 for each z and have at most a countable number of discontinuities, $\sqrt{n}(\hat{\mu} - \mu_{\text{true}})$ tends under the (2.2) assumption in distribution to

$$\Lambda = \sum_S c(S \mid D)\Lambda_S = \left(\frac{\partial \mu}{\partial \theta}\right)^t J_{00}^{-1}M + \omega^t\{\delta - \hat{\delta}(D)\}.$$

Its mean and variance are $\omega^t\{\delta - \text{E}\hat{\delta}(D)\}$ and $\tau_0^2 + \omega^t\text{Var}\,\hat{\delta}(D)\,\omega$, with mean squared error $\text{E}\Lambda^2 = \tau_0^2 + R(\delta)$, in which

$$R(\delta) = \text{E}(\omega^t\hat{\delta} - \omega^t\delta)^2 = \omega^t\,\text{E}\{\hat{\delta}(D) - \delta\}\{\hat{\delta}(D) - \delta\}^t\omega, \quad (4.3)$$

where

$$\tau_0^2 = \left(\frac{\partial\mu}{\partial\theta}\right)^t J_{00}^{-1}\frac{\partial\mu}{\partial\theta} \quad \text{and} \quad \omega = J_{10}J_{00}^{-1}\frac{\partial\mu}{\partial\theta} - \frac{\partial\mu}{\partial\gamma}.$$

Densities of various Λ's are displayed in Figure 3, illustrating, in particular, the nonnormal nature of the limit distributions.

Remark 4.1. The theorem spells out what happens to compromise estimators under the (2.2) scenario. A reviewer has wondered whether this description is adequate, if the underlying framework also allows perturbations of the θ part of the model. To investigate this issue, consider $f_n(y) = f(y, \theta_0 + \eta/\sqrt{n}, \gamma_0 + \delta/\sqrt{n})$ instead of (2.2), where $\eta = (\eta_1, \ldots, \eta_p)^t$, along with $\mu_n = \mu(\theta_0 + \eta/\sqrt{n}, \gamma_0 + \delta/\sqrt{n})$. Then Lemmas 3.1–3.3 may be generalized, leading to parallel statements involving say $\widetilde{C}_S$ and $\widetilde{D}_S$, which now also depend on η, with a consequent expression for say $\widetilde{\Lambda}_S$, the limit variable for $\sqrt{n}(\hat{\mu}_S - \mu_n)$. It turns out that $\widetilde{\Lambda}_S$ has the same distribution as before, independent of η. This shows that the description given in Theorem 4.1 continues to be adequate even when the θ part of the model is being locally perturbed.

Remark 4.2. The theorem was stated in a form focusing on D_n of (3.1) and its limit form $D \sim N_q(\delta, K)$. It is convenient, also for interpretational purposes, to rephrase in terms of

$$Z_n = \widehat{K}^{-1/2}D_n = \widehat{K}^{-1/2}\sqrt{n}(\hat{\gamma}_{\text{full}} - \gamma_0) \qquad (4.4)$$

and its limit form $Z = K^{-1/2}D \sim N_q(a, I)$, via the link $a = K^{-1/2}\delta$. Note that $Z_n \overset{d}{\to} N_q(a, I)$. Here $\widehat{K}$ is any reasonable estimator of K; it suffices that it is consistent for K under the null model $\gamma = \gamma_0$. Also, the weights of the compromise estimator $c(S \mid D_n)$ may be seen as functions of Z_n rather than of D_n. For such compromise estimators $\sum_S \bar{c}(S \mid Z_n)\hat{\mu}_S$, the limiting distribution has risk $\tau_0^2 + \overline{R}(a)$, where $\overline{R}(a) = \text{E}(\omega^t K^{1/2}\hat{a} - \omega^t K^{1/2}a)^2$ and $\hat{a}(Z) = \sum_S \bar{c}(S \mid Z)H_S Z$. This is viewed as an estimator of a on the canonical scale where $Z \sim N_q(a, I)$. With this notation $\text{AIC}_{n,S} = Z_n^t H_S Z_n - 2|S| + o_p(1)$.

Remark 4.3. Note that $\pi_S^t K_S \pi_S$ is the $q \times q$ matrix with the elements of K_S placed according to the indexes of the subset S and with 0s elsewhere. It is also worthwhile recording the simpler structure that results in the special case of a diagonal K matrix. Then H_S is diagonal with values 1 for $j \in S$ and 0 for $j \notin S$. Accordingly,

$$\hat{a}(z) = (W_1(z)z_1, \ldots, W_q(z)z_q)^t,$$

$$\text{where} \quad W_j(z) = \sum_S \bar{c}(S \mid z)I\{j \in S\} \quad (4.5)$$

in such situations. The limiting risk is $\tau_0^2 + \text{E}[\sum_{j=1}^{q}\omega_j k_j^{1/2} \times \{W_j(Z)Z_j - a_j\}]^2$. Equation (4.5) also shows that different-looking compromise strategies may well have the same performance for large n. For illustration let $q = 3$ with a diagonal K, with eight weight functions, say $\bar{c}_{000}(Z_n), \ldots, \bar{c}_{111}(Z_n)$ with 0 and 1 indicating exclusion and inclusion of γ_1, γ_2, and γ_3 in the

model. Then the performance of the procedure is determined by the three functions $W_1 = \bar{c}_{100} + \bar{c}_{101} + \bar{c}_{110} + \bar{c}_{111}$, $W_2 = \bar{c}_{010} + \bar{c}_{110} + \bar{c}_{011} + \bar{c}_{111}$, and $W_3 = \bar{c}_{001} + \bar{c}_{101} + \bar{c}_{011} + \bar{c}_{111}$.

Theorem 4.1 spells out the drastic reduction in complexity by comparing model choice and estimation strategies in the large-sample limit experiment. Performances of such regimes are characterized fully by (4.3), in other words, by a simpler estimation problem in a standard situation involving a multi-normal $D \sim N_q(\delta, K)$ with known variance matrix. Two viewpoints can be taken here. The first is that components $\delta_1, \ldots, \delta_q$ are being estimated simultaneously on the basis of D, with loss function $\{\sum_{j=1}^{q} \omega_j(\hat{\delta}_j - \delta_j)\}^2$. The alternative viewpoint is that only the one-dimensional parameter $\psi = \omega^t \delta = \omega^t K^{1/2} a$ matters and that this parameter has to be estimated under quadratic loss by estimators of the form

$$\hat{\psi} = \omega^t \hat{\delta}(D) = \omega^t K^{1/2} \left\{ \sum_S c(S \mid D) H_S \right\} K^{-1/2} D. \quad (4.6)$$

It is instructive to see the role of the parameter of interest $\mu = \mu(\theta, \gamma)$; what is in the end a good model selection strategy or a regime for smoothing between models does depend on the parameter under study. This is perhaps only to be expected, but the point is often overlooked, in that the most popular model choice methods work independently of the inference to take place afterwards. See Claeskens and Hjort (2003) for applications where different estimands correspond to different optimal submodels.

Importantly, the theory developed previously goes through also for the regression model cases, under mild regularity conditions of the type described in Section 3.4.

4.2 Dwindling Confidence

The traditional use of model selection methods in practice is to proceed as if the finally selected model had been chosen a priori. Thus, a typical confidence interval, taking intended coverage probability 95% as an example, would take the form

$$\mu \in \hat{\mu}_{\widehat{S}} \pm 1.96 \, \hat{\tau}_{\widehat{S}}/\sqrt{n}, \quad (4.7)$$

where $\widehat{S}$ represents the chosen model and $\hat{\tau}_S/\sqrt{n}$ is an estimator of the standard deviation for $\hat{\mu}_S$, without model uncertainty for S. From Lemma 3.3 and Theorem 4.1, $\hat{\tau}_S$ estimates $\tau_S = (\tau_0^2 + \omega^t K^{1/2} H_S K^{1/2} \omega)^{1/2}$. Such procedures ignore the uncertainties involved in the model selection step of the analysis and are, consequently, too optimistic about the confidence level attained by such intervals; similar comments apply to tests and other forms of inference. This is, for example, visible when one compares the optimistic standard deviation estimates of Table 1, for the AIC-chosen μ estimators, with the real ones, as found in Table 2, column 6.

Consider any selection estimator of this type, where the model selection is being determined exactly or asymptotically via Z_n of (4.4). These correspond to compromise estimators (4.1) for which $\mathcal{R}^q$ is partitioned into regions R_S, where $c(S \mid z) = 1$ for $z \in R_S$ and 0 outside. Now study

$$V_n = \sqrt{n}(\hat{\mu}_{\widehat{S}} - \mu_{\text{true}})/\hat{\tau}_{\widehat{S}}.$$

By previous efforts $V_n \xrightarrow{d} V = \Lambda/\tau(Z)$, say, where $\tau(z)^2 = \tau_0^2 + \omega^t K^{1/2} H_S K^{1/2} \omega$ for $z \in R_S$. Also, from the proof of


Figure 1. True Coverage Probability When Ignoring AIC Choice Among Four Models for $q = 2$ When $\omega = (1, 1)^t$ and $K = diag(1, 1)$.

Theorem 4.1, $\Lambda \mid z$ is normal with variance τ_0^2 and mean $\omega^t \{\delta - G(z)^t K^{1/2} z\}$. Thus, the real coverage probability of an interval like (4.7) goes for growing n to

$$p(a) = \text{Pr}_a\{|V| \leq 1.96\}$$

$$= \sum_S \int_{R_S} \text{Pr}\{|E(\Lambda \mid z) + \tau_0 N|/\tau_S \leq 1.96\}\phi(z - a)\, dz,$$

where N denotes a standard normal variable and again $a = K^{-1/2}\delta$.

For $q = 1$ these probabilities are easily calculated via numerical integration. For the AIC-selected estimator, one chooses the narrow estimator when $|Z| \leq \sqrt{2}$ and the full one when $|Z| > \sqrt{2}$, and some algebra leads to $p(a)$ being equal to

$$\int_{|z| \leq \sqrt{2}} \text{Pr}\{|\rho a + N| \leq 1.96\}\phi(z - a)\, dz$$

$$+ \int_{|z| > \sqrt{2}} \text{Pr}\left\{ \frac{|\rho(z - a) + N|}{(1 + \rho^2)^{1/2}} \leq 1.96 \right\}\phi(z - a)\, dz,$$

in terms of $\rho = \omega K^{1/2}/\tau_0$. This is often significantly smaller than the intended level .95. Figure 2 displays the true coverage probability as a function of a, for AIC model choice between narrow and wide models. We have also carried out such computations for the case of $q = 2$, using simulations. Figure 1 presents the coverage deficiency for AIC choice among four models in a situation where $\omega = (1, 1)^t$ and $K = \text{diag}(1, 1)$. Correct coverage is obtained in the limit as $\|a\| \to \infty$.

4.3 Better Confidence

We have seen that the traditionally employed construction (4.7) leads to too optimistic intervals, in that the real coverage probability is lower than the intended level. Aware of this phenomenon, Buckland et al. (1997) suggested a method for taking the extra model uncertainty into account, which, in particular, leads to modified confidence intervals. Their method was later embraced by Burnham and Anderson (2002, sec. 4.3), particularly in conjunction with the smoothed AIC weights for $c(S \mid D_n)$; see Section 5.2. The method amounts to using $\hat{\mu} \pm u \, \widehat{\text{se}}_n$ as confidence intervals, with u the appropriate normal

quantile and formula (9) in Buckland et al. for the estimated standard error $\widehat{\text{se}}_n$. Rephrased to fit our framework,

$$\widehat{\text{se}}_n = \sum_S c(S \mid D_n)(\hat{\tau}_S^2/n + \hat{b}_S^2)^{1/2},$$

where $\hat{\tau}_S$ is a consistent estimator of $\tau_S = (\tau_0^2 + \omega^t K^{1/2} H_S K^{1/2} \omega)^{1/2}$ and $\hat{b}_S = \hat{\mu}_S - \hat{\mu}$. The resulting coverage probability p_n is not studied accurately in the references mentioned, but it is claimed that it will be close to the intended $\Pr\{-u \leq \text{N}(0, 1) \leq u\}$. Our methods make it possible to study p_n accurately, however. One has $p_n = \Pr\{-u \leq B_n \leq u\}$, where $B_n = (\hat{\mu} - \mu_{\text{true}})/\widehat{\text{se}}_n$. This variable has a well-defined limit distribution, because $\sqrt{n}\widehat{\text{se}}_n \xrightarrow{d} \widehat{\text{se}} = \sum_S c(S \mid D)\{\tau_S^2 + (\Lambda_S - \Lambda)^2\}^{1/2}$, simultaneously with $\sqrt{n}(\hat{\mu} - \mu_{\text{true}}) \xrightarrow{d} \Lambda$, by an extension of arguments used in the Appendix to prove Theorem 4.1. Furthermore, $\Lambda_S - \Lambda = \omega^t\{\hat{\delta}(D) - K^{1/2} H_S K^{-1/2} D\}$. Thus,

$$B_n \xrightarrow{d} B = \frac{\Lambda}{\widehat{\text{se}}}$$

$$= \frac{\Lambda_0 + \omega^t\{\delta - \hat{\delta}(D)\}}{\sum_S c(S \mid D)\{\tau_S^2 + [\omega^t\{\hat{\delta}(D) - K^{1/2} H_S K^{-1/2} D\}]^2\}^{1/2}},$$

writing $\Lambda_0 = (\partial\mu/\partial\theta)^t J_{00}^{-1} M$. The B variable is a normal, for given D, but is clearly not standard normal when averaged over the distribution of D, and neither is it centred at 0, so the coverage probability p_n is biased.

To illustrate this, consider the $q = 1$ case, with compromise estimator $\hat{\mu} = \{1 - W(Z_n)\}\hat{\mu}_{\text{narr}} + W(Z_n)\hat{\mu}_{\text{full}}$, for which $\Lambda_\varnothing = \Lambda_0 + \omega K^{1/2} a$ and $\Lambda_{\text{full}} = \Lambda_0 + \omega K^{1/2}(a - Z)$. Here $B = \Lambda/\widehat{\text{se}}$ takes the form

$$\big(\Lambda_0 + \omega K^{1/2}\{a - W(Z)Z\}\big)\big/$$

$$\big(\{1 - W(Z)\}\{\tau_0^2 + \omega^2 K W(Z)^2 Z^2\}^{1/2}$$

$$+ W(Z)\{\tau_0^2 + \omega^2 K + \omega^2 K\{1 - W(Z)\}^2 Z^2\}^{1/2}\big)$$

with $\Lambda_0 \sim \text{N}(0, \tau_0^2)$ and independent of $Z \sim \text{N}(a, 1)$. The limiting coverage probability may then be computed, via numerical integration, as $p(a) = \int \Pr\{-u \leq B \leq u \mid z\}\phi(z - a)\,dz$. See Figure 2.

Consider instead

$$\begin{aligned} \text{low}_n &= \hat{\mu} - \widehat{\omega}^t\{D_n - \hat{\delta}(D_n)\}/\sqrt{n} - u\hat{\kappa}/\sqrt{n}, \\ \text{up}_n &= \hat{\mu} - \widehat{\omega}^t\{D_n - \hat{\delta}(D_n)\}/\sqrt{n} + u\hat{\kappa}/\sqrt{n}, \end{aligned} \quad (4.8)$$

where $\widehat{\omega}$ and $\hat{\kappa}$ are consistent estimators of ω and $\kappa = \tau_{\text{full}} = (\tau_0^2 + \omega^t K\omega)^{1/2}$ and u is a normal quantile. We observe that the coverage probability $p_n = \Pr\{\text{low}_n \leq \mu_{\text{true}} \leq \text{up}_n\}$ is the same as $\Pr\{-u \leq T_n \leq u\}$, where

$$T_n = \big[\sqrt{n}(\hat{\mu} - \mu_{\text{true}}) - \widehat{\omega}^t\{D_n - \hat{\delta}(D_n)\}\big]/\hat{\kappa}.$$

But there is simultaneous convergence in distribution

$$\big(\sqrt{n}(\hat{\mu} - \mu_{\text{true}}), D_n\big) \xrightarrow{d} \big(\Lambda_0 + \omega^t\{\delta - \hat{\delta}(D)\}, D\big),$$

essentially by the arguments used to prove Theorem 4.1. It follows that $T_n \xrightarrow{d} \{\Lambda_0 + \omega^t(\delta - D)\}/\kappa$, which is simply a standard normal. Thus, with $u = 1.645$, for example, the (4.8) interval has asymptotic confidence level precisely the intended 90% level.


Figure 2. Exact Limiting Coverage Probability p(a) for Three Confidence Interval Procedures in Two Situations, Corresponding to $\rho = \omega K^{1/2}/\tau_0$ Equal to 1 and 2/3, for q = 1. The three methods are the AIC-based version of (4.7) (dotted line); the smoothed AIC method of Section 5.2 using $\widehat{\text{se}}_n$ described previously as standard error (solid line); and, finally, the general (4.8) method, which gives correct .90 coverage for each method (dashed line).

4.4 Example: Exponential Within Weibull

Let $Y_1, \ldots, Y_n$ come from the Weibull distribution with cumulative $1 - \exp\{-(\theta y)^\gamma\}$, with γ in the vicinity of $\gamma_0 = 1$. With some effort one finds the information matrix, with inverse:

$$J = \begin{pmatrix} \gamma^2/\theta^2 & (1-r)\theta \\ (1-r)/\theta & c^2/\gamma^2 \end{pmatrix},$$

$$J^{-1} = \frac{1}{\pi^2/6}\begin{pmatrix} c^2\theta^2/\gamma^2 & -(1-r)\theta \\ -(1-r)/\theta & \gamma^2 \end{pmatrix},$$

where $r = .5772\ldots$ is the Euler–Mascheroni constant and $c^2 = \pi^2/6 + (1 - r)^2$. We consider estimators of the median $\mu = (\log 2)^{1/\gamma}/\theta$ of the form

$$\begin{aligned} \hat{\mu} &= \{1 - W(Z_n)\}\hat{\mu}_{\text{narr}} + W(Z_n)\hat{\mu}_{\text{full}} \\ &= \{1 - W(Z_n)\}\frac{\log 2}{\hat{\theta}_{\text{narr}}} + W(Z_n)\frac{(\log 2)^{1/\hat{\gamma}_{\text{full}}}}{\hat{\theta}_{\text{full}}}, \end{aligned}$$

where, following our recipe, $Z_n = \sqrt{n}(\hat{\gamma}_{\text{full}} - 1)/\widehat{K}^{1/2}$ with $\widehat{K}$ estimating $K = 6\gamma^2/\pi^2$. Also, $\omega = v/\theta\{-(1 - r) + \log v\}$, in terms of $v = \log 2$, and we find

$$\tau_0 = v/\theta,$$

$$(K\omega^2)^{1/2} = (v/\theta)| -(1 - r) + \log v|\sqrt{6}/\pi.$$

When $\gamma = 1 + \delta/\sqrt{n}$ the limit distribution of $\sqrt{n}(\hat{\mu} - \mu_{\text{true}})$ is $\Lambda = \Lambda_0 + \omega K^{1/2}\{a - W(Z)\}$, where $\Lambda_0 \sim \text{N}(0, \tau_0^2)$ and is independent of $Z \sim \text{N}(a, 1)$, and $a = \delta/K^{1/2}$.

We have carried out simulations in this example, for estimation of the median and other quantiles, using hard and smoothed AIC estimators, and yet further of the compromise estimators described in Section 5. The density of T_n was seen to be quite close to its limiting standard normal density, for even moderate n. The coverage probability for the (4.8) intervals is, consequently, close to the intended level.

5. SOME MODEL AVERAGE ESTIMATION SCHEMES

In this section we go through a partial list of particularly attractive FMA methods. Different FMA schemes are characterized by their δ-estimator and ψ-estimator counterparts $G(D)^{\mathrm{t}}D$ and $\omega^{\mathrm{t}}G(D)^{\mathrm{t}}D$ in the limit experiment, as shown in the previous section. It is therefore often fruitful to construct FMA regimes via arguments inside the context of the limit experiment.

5.1 The AIC Selection Estimator

For the AIC method with all 2^q subsets allowed, let R_S be the set of D such that $\mathrm{AIC}_S(D)$ is larger than all other $\mathrm{AIC}_{S'}(D)$, where

$$\mathrm{AIC}_S(D) = D^{\mathrm{t}}K^{-1/2}H_SK^{-1/2}D - 2|S| = Z^{\mathrm{t}}H_SZ - 2|S|. \quad (5.1)$$

Then, for $D \in R_S$, $c(S \mid D) = 1$ whereas the other $c(S' \mid D) = 0$. For the case of K being a diagonal matrix with diagonal elements k_j, we have $\mathrm{AIC}_S(D) = \sum_{j \in S}(D_j^2/k_j - 2)$. This shows that, to the first order of large-sample approximation, precisely those j are included in the selected set for which $D_{n,j}^2/\hat{k}_j = n(\hat{\gamma}_{\mathrm{full},j} - \gamma_{0,j})^2/\hat{k}_j > 2$.

5.2 A Smoothed AIC-Based Estimator

Buckland et al. (1997) made a general model averaging suggestion that amounts to taking weights $c^*(S \mid \mathrm{data})$ proportional to $\exp(\ell_S - |S|)$, where ℓ_S is the maximized log-likelihood at model S. Thus, comparing weights for models of the same complexity corresponds to likelihood ratio methods, and the penalization of these otherwise ad hoc constructed terms stems from the analogy with the AIC method. By (3.3) the smoothed AIC weights may be represented as

$$\frac{\exp(\frac{1}{2}\mathrm{AIC}_{n,S})}{\sum_{\mathrm{all}\ S'}\exp(\frac{1}{2}\mathrm{AIC}_{n,S'})}$$

$$= \frac{\exp(\frac{1}{2}D_n^{\mathrm{t}}K^{-1/2}H_SK^{-1/2}D_n - |S|)}{\sum_{\mathrm{all}\ S'}\exp(\frac{1}{2}D_n^{\mathrm{t}}K^{-1/2}H_S'K^{-1/2}D_n - |S'|)} + o_p(1). \quad (5.2)$$

It follows from the theory developed in Section 4 that the large-sample distributions of compromise estimators are the same, whether one uses the left-hand-side ratio or the right-hand-side ratio as weights. Note also that there is some independent motivation for using such weights from a Bayesian analogy, where $\exp(\frac{1}{2}\mathrm{BIC}_S)/\sum_{S'}\exp(\frac{1}{2}\mathrm{BIC}_{S'})$ is known to be an approximation to the posterior probability of model S being correct; see Schwarz (1978), as well as the discussion in Burnham and Anderson (2002, sec. 6.4). Results developed in Section 8 lead to other, potentially better approximations.

Using the theory developed in Section 4, the limiting distribution is a suitable convex mixture of normals, and the limiting squared error can be computed via (4.3). Buckland et al. partly motivated their method by considering correlations between different estimators, but without estimating these correlations accurately. We may show, using arguments of Section 4, that the limiting correlation between submodel estimators $\hat{\mu}_S$ and $\hat{\mu}_{S'}$ is

$$\rho(S, S') = (\tau_0^2 + \omega^{\mathrm{t}}K^{1/2}H_SH_{S'}K^{1/2}\omega)$$
$$\times (\tau_0^2 + \omega^{\mathrm{t}}K^{1/2}H_SK^{1/2}\omega)^{-1/2}$$
$$\times (\tau_0^2 + \omega^{\mathrm{t}}K^{1/2}H_{S'}K^{1/2}\omega)^{-1/2}. \quad (5.3)$$

We may also derive the limiting correlation between any two compromise estimators via similar arguments. Its size depends, in particular, on the relative sizes of τ_0 and $(\omega^{\mathrm{t}}K\omega)^{1/2}$.

5.3 The FIC Selection Estimator

The AIC method selects one winning model, regardless of the intended use for this model. In contrast, Claeskens and Hjort (2003) developed a focused information criterion that specifically aims at finding the best candidate model for a given focus parameter μ. Whereas the AIC method chooses S to maximize $\mathrm{AIC}_S(D)$ of (5.1), the FIC goes for S to minimize

$$\mathrm{FIC}_S(D) = (\omega^{\mathrm{t}}D - \hat{\psi}_S)^2 + 2\omega_S^{\mathrm{t}}K_S\omega_S,$$

$$\text{where} \quad \hat{\psi}_S = \omega^{\mathrm{t}}K^{1/2}H_SK^{-1/2}D.$$

This is the limit experiment version of the FIC. In practice, one plugs in estimates of ω, K, K_S, and H_S. Suppose, for example, that the choice is only between the narrow and the full models. Then the AIC selects the full model provided $D^{\mathrm{t}}K^{-1}D \geq 2q$, whereas the corresponding FIC criterion for selecting the full model is $(\omega^{\mathrm{t}}D)^2 \geq 2\omega^{\mathrm{t}}K\omega$.

It is also attractive to smooth across estimators using the information carried by the FIC scores, and we suggest using

$$c(S \mid D) = \exp\left(-\frac{1}{2}\kappa\frac{\mathrm{FIC}_S}{\omega^{\mathrm{t}}K\omega}\right)\Bigg/$$

$$\sum_{\mathrm{all}\ S'}\exp\left(-\frac{1}{2}\kappa\frac{\mathrm{FIC}_{S'}}{\omega^{\mathrm{t}}K\omega}\right), \quad \text{with} \quad \kappa \geq 0. \quad (5.4)$$

Here κ is an algorithmic parameter, bridging from uniform weighting (κ close to 0) to the hard-core FIC (which is the case of large κ). Of course, the added $\frac{1}{2}$ is somewhat redundant, but the form (5.4) is suggested by connections to certain empirical Bayes arguments that can be developed using the theory of Section 9. The $\omega^{\mathrm{t}}K\omega$ factor appearing in the scaling for κ is the constant risk of the minimax estimator $\hat{\delta} = D$. The point of the scaling is that κ values used in different data contexts can now be compared directly. One may show here that, for the one-dimensional case $q = 1$, the value $\kappa = 1$ makes the weights of (5.4) agree with those for the smoothed AIC.

To illustrate the limiting distribution of some compromise estimators, we display in Figure 3 the density of Λ for a situation with $q = 2$ extra parameters, where $K = \mathrm{diag}(1, 1)$, $\omega = (1, 1)^{\mathrm{t}}$, and $\tau_0 = 0.5$. The nonnormal nature is evident, not only for nonsmooth methods such as the AIC, but also for smoothed versions thereof. For each of the four positions in the parameter space considered here, the smoothed FIC wins over the others in terms of mean squared error.

5.4 Minimizing Estimated Risk

Consider estimators of the form $\sum_S c(S)\hat{\mu}_S$, with nonrandom weights summing to 1. From previous results the limiting distribution in question is that of $\Lambda = \sum_S c(S)\Lambda_S$, with Λ_S as in Lemma 3.3. One finds $E\Lambda = \omega^{\mathrm{t}}(I - Q)^{\mathrm{t}}\delta$, where $Q = \sum_S c(S)K^{-1/2}H_SK^{1/2}$; furthermore, $\mathrm{Var}\,\Lambda = \tau_0^2 + \omega^{\mathrm{t}}Q^{\mathrm{t}}KQ\omega$, using the covariance extension of Lemma 3.3, which was also used in connection with (5.3). Thus, the limiting risk of the estimator is $\tau_0^2 + R(\delta)$, where $R(\delta) = \omega^{\mathrm{t}}\{(I - Q)^{\mathrm{t}}\delta\delta^{\mathrm{t}}(I - Q) + Q^{\mathrm{t}}KQ\}\omega$. This also agrees with (4.3).

Figure 3. Density of the Limiting Distribution Λ of $\sqrt{n}(\hat{\mu} - \mu_{true})$ for Three Compromise Estimators at Four Positions in the Parameter Space. The situation studied has $q = 2$, $K = diag(1, 1)$, $diff = (1, 1)^t$, and $\tau_0 = .5$, and the four positions are (a) (0,0), (b) (1.5, 1.5), (c) (1, −1), and (d) (2, −2) for $a = (a_1, a_2)$. The estimators are post-AIC (solid line), smoothed AIC (dotted line), and smoothed FIC with $\kappa = 1$ in (5.4) (dashed line).

Various model average estimators may now be constructed along the following lines. Estimate the risk $R(\delta)$, for example, by inserting D for δ, or, alternatively, the unbiased $DD^t - K$ for $\delta\delta^t$. Then select weights $c(S)$ to minimize this estimated risk. Different versions emerge from this, depending also on the list of submodels one wishes to smooth across. A simple special case worth recording is that of smoothing optimally between the two extreme models, $\hat{\mu} = (1 - c)\hat{\mu}_{narr} + c\hat{\mu}_{full}$. Using $\hat{\psi} = \omega^t D$ as an estimator of $\psi = \omega^t \delta$ in the full model, for the current purpose of estimating the optimal weights, the result for the limit experiment situation is

$$\hat{\mu} = \frac{\omega^t K \omega}{\hat{\psi}^2 + \omega^t K \omega}\hat{\mu}_{narr} + \frac{\hat{\psi}^2}{\hat{\psi}^2 + \omega^t K \omega}\hat{\mu}_{full}. \quad (5.5)$$

With real data one, in addition, plugs in estimates of ω and K and uses $\hat{\psi} = \widehat{\omega}^t D_n$.

5.5 Smoothing Across Singletons

An attractive challenge is to form data-based averages over estimators $\hat{\mu}_{\{j\}}$, corresponding to the simple one-parameter model extensions of the narrow model. These take the form $\sum_{j=0}^{q} c(j \mid D_n)\hat{\mu}_{\{j\}}$, where $j = 0$ corresponds to the narrow model estimator and might be thought of as resembling first-stage Taylor expansion estimators. These μ estimators are further related to δ estimators of the form $\hat{\delta} = G(D)^t D$ in the limit experiment, where $G(D)$ is as in (4.2) but engaging only H_S matrices for S being empty or a singleton.

For brevity we present only one of these methods, which has been shown to perform well in some limited simulation exercises of the authors. This method emerges from Bayesian and

empirical Bayesian considerations, starting with a prior that has $\delta = 0$ with some probability p_0 and with probability p_j has δ_j from a normal and the other δ_i's equal to 0, and where $\sum_{j=0}^{q} p_j = 1$. The estimator is

$$\hat{\mu} = (1 - \hat{\rho})\hat{\mu}_{narr} + \hat{\rho} \sum_{j=1}^{q} \frac{\exp(\frac{1}{2}\hat{\rho}k^{jj}\widehat{T}_j^2)}{\sum_{i=1}^{q} \exp(\frac{1}{2}\hat{\rho}k^{ii}\widehat{T}_i^2)}\hat{\mu}_{\{j\}},$$

$$\text{where} \quad \hat{\rho} = \frac{\hat{\tau}^2}{1 + \hat{\tau}^2}, \quad (5.6)$$

with $\hat{\tau} = (D_n^t \widehat{K}^{-1} D_n - q)_{+}^{1/2}$ and $\widehat{T}_j = (\hat{k}^{jj})^{-1}e_j^t\widehat{K}^{-1}\hat{\delta}_{full}$, in terms of the diagonal elements of $\widehat{K}^{-1}$ and the jth unit vector e_j. Details of this construction, along with useful variations, are available in a technical report from the authors.

5.6 An Empirical Bayes Model Smoother

The following arguments motivate a particular estimator smoother, with data-dependent weights $c(S \mid D_n)$ in (4.1). The idea is to start with a Bayesian mixture prior, of a more general type than that used in Section 5.5, then work out the necessary details pertaining to the posterior, and, finally, estimate the required spread parameter from the marginal distribution of data.

Remark 5.1. We take this opportunity to make the following general point. Our theory has been developed by the desire to handle averages of subset estimators $\hat{\mu}_S$ of the form (2.1), that is, for subsets of the original (θ, γ) or (θ, δ) parameterization of the fullest model. Mathematically, we are free to reparameterize from δ to the canonical $a = K^{-1/2}\delta$ scale, however, and instead

work with subset estimators $\mu_S^* = \mu(\theta_S^*, a_S^*, 0_{S^c})$ and averages $\mu^* = \sum_S c^*(S \mid Z_n)\mu_S^*$. The advantage is a cleaner orthogonal structure, because $Z \sim N_q(a, I)$. The theory of Sections 3 and 4 would go through with minor changes. We illustrate this here, because the mixture strategy becomes easier to develop and describe.

Focus first on one of the a_j components, and let it be 0 with probability p_0 and an $N(0, \sigma^2)$ with probability p_1. Then $a_j \mid z_j$ is 0 with probability $\tilde{p}_0(z_j)$ and from an $N(\rho z_j, \rho)$ with probability $\tilde{p}_1(z_j)$, where $\rho = \sigma^2/(1 + \sigma^2)$. Furthermore,

$$\tilde{p}_1(z_j) = \frac{p_1 \phi(z_j, 1 + \sigma^2)}{p_0 \phi(z_j, 1) + p_1 \phi(z_j, 1 + \sigma^2)}$$

$$= \frac{p_1 (1 + \sigma^2)^{-1/2} \exp(\frac{1}{2}\rho z_j^2)}{p_0 + p_1(1 + \sigma^2)^{-1/2} \exp(\frac{1}{2}\rho z_j^2)},$$

with $\tilde{p}_0(z_j) = 1 - \tilde{p}_1(z_j)$ and $\phi(z, v^2)$ the $N(0, v^2)$ density evaluated at z. If now $a_1, \ldots, a_q$ are given independent priors of this type, which is reasonable in that the a_j's have been transformed toward orthogonality and the same scale, then $E(a_j \mid z) = \rho \tilde{p}_1(z_j)z_j$ for $j = 1, \ldots, q$. By the general recipe established in Remark 4.3, we should have $\hat{a}_j = W_j(z)z_j$ for $W_j(z) = \sum_{S: j \in S} c^*(S \mid z)$. But this fits in with the compromise regime that uses

$$c^*(S \mid z) = \rho \prod_{j=1}^q \tilde{p}_0(z_j)^{I\{j \notin S\}} \tilde{p}_1(z_j)^{I\{j \in S\}} \quad \text{for nonempty } S$$

and $c(\varnothing \mid z) = 1 - \rho + \rho \prod_{j=1}^q \tilde{p}_0(z_j)$. A fruitful variation is the empirical Bayes construction, which inserts an estimate $\hat{\sigma}$ for σ in the $c^*(S \mid z)$ formulas given previously. Such an estimate may emerge from likelihood analysis based on the marginal distribution of $(Z_1, \ldots, Z_q)$. One may also use a hyperprior for σ in a two-stage Bayesian fashion. It suffices for the present purposes to devise a simple moment estimator, however, using that $\sum_{j=1}^q Z_j^2$ has mean $q + qp_1\sigma^2$. We therefore propose $\hat{\sigma}^2 = (\sum_{j=1}^q Z_j^2 - q)_+/(qp_1)$, where the positive part notation indicates that $\hat{\sigma} = 0$ in the case of $\sum_{j=1}^q Z_j^2 \leq q$. Such an event suggests that none of the a_j's are significantly nonzero, and the scheme selects the narrow model.

Several variations of these arguments could be considered. For example, one may use a vague hyperprior for the σ parameter. Another alternative is to estimate both σ and p_1 in the preceding construction based on the marginal distribution of $(Z_1, \ldots, Z_q)$, which obviates the need to specify p_1 in advance.

6. ILLUSTRATIONS AND APPLICATIONS

6.1 Computational Aspects

Frequentist model averaging analysis can be easily performed using standard statistical software. All numerical results presented in this article are obtained using the software packages S-Plus and R.

Obtain parameter estimates in the different models and either compute the model selection criterion value for each of these models in order to form the indicator variable of the optimal model or directly construct the general model averaging weights of choice. From the estimate in the biggest model we

construct $\hat{\delta}$. Nonlinear optimization algorithms, such as `nlm()` in R, provide us immediately with a matrix of second-order partial derivatives, leading to the matrix $\hat{J}_{\text{full}}$. Next we construct the projection matrices π_S and use these to define $\hat{K}_S$ and $\hat{H}_S$ for each model S. Partial derivatives of μ wrt θ parameters and γ parameters, at either $(\hat{\theta}_{\text{narr}}, \gamma_0)$ or $(\hat{\theta}_{\text{full}}, \gamma_0)$, are needed for the computation of $\hat{\omega}$ and $\hat{\tau}_0$. These are sometimes easy to derive mathematically and can otherwise be computed using numerical derivatives.

As far as our theoretical results are concerned, we may use any J_{full}^* estimator for the crucial matrix J_{full} of Section 3.1 (along with the consequent estimators for K, K_S, and H_S), as long as it is consistent under our $\gamma_0 + \delta/\sqrt{n}$ framework. In particular, it may be computed under "narrow" or "full" circumstances. Narrow estimation is sometimes easiest, via explicit formulas or via simulation of score vectors under the null model. To guard against cases where δ is some distance away from 0, however, it will be more satisfactory and robust to use full-model estimation; see also a parallel discussion of this in Claeskens and Hjort (2003).

We have found it useful in practice to simulate the limit distribution Λ of Theorem 4.1 for the average estimator scheme being used at δ corresponding to its estimate $\hat{\delta}_{\text{full}}$. A density estimate of say 10,000 such Λ copies is informative and leads to estimated bias and standard deviation for the compromise estimator being used, as well as to approximative confidence intervals.

6.2 Averaging Over Logistic Regression Models

Time has come to revisit the 189 babies of Section 1.2. Here we illustrate our general methodology by exhibiting results for each of the three focus parameters $p(\text{white})$, $p(\text{black})$, and their ratio $p(\text{black})/p(\text{white})$ for six different FMA regimes (see Table 2). These are the AIC-selected estimator; the smooth AIC of Section 5.2; the FIC-selected estimator; the smooth FIC of (5.4) with $\kappa = 1$; the smoothing across singletons, which makes data-dictated compromises among the four models "0", "3", "4", and "5"; and, finally, the simple compromise between narrow and full models as in (5.5).

We record here that, for $p(\text{white})$, $\hat{\omega} = (-.245, .032, .065)$ and $\hat{\tau}_0 = .477$; for $p(\text{black})$, $\hat{\omega} = (.429, -.185, .073)$ and $\hat{\tau}_0 = .550$; whereas, for the ratio parameter, $\hat{\omega} = (3.783, -10.057, -.190)$ with $\hat{\tau}_0 = .495$. We observe that the averaging across singletons method leads to low standard deviation and short confidence intervals. This is particularly noticeable for the ratio parameter. The standard deviations and the confidence bounds come from 10,000 simulations of the appropriate Λ distributions.

6.3 Averaging Over Covariance Structure Models

There are no inherent problems with applying our methodology to situations where data are multidimensional. To illustrate this, we report on a brief investigation of multinormal data where different models for the covariance structure, in the absence of clear a priori preferences, are being averaged over to form estimators of quantities of interest. We note that there are several areas of statistics where covariance modeling is of interest, and sometimes perhaps of primary concern, as with factor analysis, and where variations of our methods might be fruitful.

Table 2. For Each of the Three Focus Parameters Associated With the Study of Low Birth Weights Described in Section 1.2, the Table Gives Parameter Estimate and Estimated Standard Deviation, Along With Lower and Upper Points for 90% Confidence Intervals for the FMA Strategies

	(a)	(b)	(c)	(d)	(e)	(f)
For p(white)						
Estimate	.269	.261	.263	.258	.281	.242
Standard deviation	.051	.047	.050	.045	.039	.048
Lower	.174	.168	.173	.165	.173	.191
Upper	.343	.322	.338	.315	.302	.350
For p(black)						
Estimate	.412	.368	.412	.365	.323	.380
Standard deviation	.112	.107	.115	.106	.107	.096
Lower	.257	.216	.257	.215	.203	.190
Upper	.618	.559	.614	.553	.549	.501
For the ratio p(black)/p(white)						
Estimate	1.533	1.440	1.564	1.501	1.159	1.651
Standard deviation	.668	.610	.647	.582	.516	.567
Lower	.681	.640	.712	.779	.843	.384
Upper	2.792	2.563	2.758	2.619	2.414	2.299

NOTE: (a) AIC, (b) smooth-AIC, (c) FIC, (d) smooth-FIC, (e) smoothing across singletons, and (f) compromise between narrow and full model.

Assume one has observed d-dimensional vectors $Y = (X_1, \ldots, X_d)^t$ from the multinormal $N_d(\xi, \Sigma)$, where different models for the structure of Σ are being considered. As a specific example, we use data from the so-called Adelskalenderen of speedskating. This is the list of the best speedskaters ever, as ranked by their personal bests over the four distances 500, 1,500, 5,000, and 10,000 m, via the classical point sum $X_1 + X_2 + X_3 + X_4$, where X_1 is the 500-m time, X_2 is the 1,500-m time divided by 3, X_3 is the 5-k time divided by 10, and X_4 the 10-k time divided by 20. The correlation structure of the 4-vector Y is important when relating, discussing, and predicting performances on different distances. Whereas there is a long list of parameters $\mu = \mu(\xi, \Sigma)$ that on occasions will ignite the fascination of speedskating fans, for this discussion we single out as focus parameters the generalized standard deviation measures $\mu_1 = \{\det(\Sigma)\}^{1/8}$ and $\mu_2 = \{\mathrm{Tr}(\Sigma)\}^{1/2}$, the average correlation $\mu_3 = \frac{1}{6} \sum_{i<j} \mathrm{corr}(X_i, X_j)$, and the so-called maximal correlation μ_4 between (X_1, X_2, X_3) and X_4. The latter is the maximal correlation between a linear combination of X_1, X_2, X_3, and X_4, and is, for example, of interest at championships when one tries to predict the final outcomes, after the completion of the three first distances. It is also equal to $(\Sigma_{10} \Sigma_{00}^{-1} \Sigma_{01} / \Sigma_{11})^{1/2}$, in terms of the blocks of Σ, of size 3×3 for Σ_{00} and so on. Below we analyze the top of the Adelskalenderen, with the best $n = 250$ skaters ever, as per the end of the 2002 season. The vectors $Y_1, \ldots, Y_n$ are, by definition, ranked, but as long as one discusses estimators that are permutation invariant we may view the data vectors as a random sample from the population of the top skaters of the world.

A minimal plausible model for Σ is M_0, which has equicorrelation and equal variances. Model M_1 assumes equicorrelation although it allows the variances to be different, whereas M_2 has no preimposed structure on the correlations but does assume equal variances. The fullest model M_3 is the unstructured covariance matrix with 10 parameters. To place this setting into the framework developed in earlier sections, let

$$\Sigma = \sigma^2 \begin{pmatrix} 1 & \rho\phi_2 & \rho\phi_3(1+v_{13}) & \rho\phi_4(1+v_{14}) \\ \rho\phi_2 & \phi_2^2 & \rho\phi_2\phi_3(1+v_{23}) & \rho\phi_2\phi_4(1+v_{24}) \\ \rho\phi_3(1+v_{13}) & \rho\phi_2\phi_3(1+v_{23}) & \phi_3^2 & \rho\phi_3\phi_4(1+v_{34}) \\ \rho\phi_4(1+v_{14}) & \rho\phi_2\phi_4(1+v_{24}) & \rho\phi_3\phi_4(1+v_{34}) & \phi_4^2 \end{pmatrix}.$$

The parameter $\theta = (\sigma^2, \rho)$ is present in all of the models, whereas subsets of $\gamma = (\phi_2, \phi_3, \phi_4, v_{13}, v_{14}, v_{23}, v_{24}, v_{34})$ are present in some of the models. Here $\gamma_0 = (1, 1, 1, 0, 0, 0, 0, 0)$. We use the criteria AIC and FIC to select an appropriate covariance structure.

For the models described previously, we get the parameter estimates shown in Table 3. Note that FIC depends on the parameter under focus and, hence, gives different values for different μ_k's. On this occasion the FIC for parameter μ_2 points to model M_1, whereas FIC selects model M_3 for all other parameters, as does the AIC. Also presented in the table are the model-averaged estimates using smoothed AIC and FIC weights, using weights as with (5.2) and (5.4), where for the latter $\kappa = 1$. Confidence intervals are constructed using (4.8) with the observed value of $D_n = \hat{\delta}_{\mathrm{full}}$ equal to $(-1.624, -.308, 5.948, -16.482, -22.861, -7.306, -15.478, -.892)^t$. At a nominal level of 90%, we find for the μ_2 parameter $(2.390, 2.611)$ for FIC, $(2.254, 2.474)$ for both AIC and smooth AIC, and $(2.339, 2.559)$ for smooth FIC.

Following the computational steps in Section 6.1, we use simulation to compute the standard deviation of the estimators for post-model selection estimation by AIC and FIC, as well as for the smoothed versions. For the μ_2 parameter, for example, we have the following estimated standard deviations for the different methods: the same value 1.029 for post-AIC and smooth AIC, whereas for post-FIC and smooth FIC the value is 1.009.

6.4 Variable Selection and Model Smoothing in Linear Regression

Assume that observations Y_i are to be regressed wrt regressors $x_{i,1}, \ldots, x_{i,p}$ and possibly wrt a further subset of additional

Table 3. Six Different Estimates of the Parameters μ_1, μ_2, μ_3, and μ_4. These correspond to models M_0, M_1, M_2, and M_3, and to AIC-smoothed and FIC-smoothed averages thereof, as per Sections 5.2 and 5.3

Parameter	M_0	M_1	M_2	M_3	Smooth AIC	Smooth FIC
μ_1	1.146	1.101	.844	.816	.816	.816
μ_2	2.364	2.389	2.461	2.364	2.364	2.381
μ_3	.225	.271	.388	.262	.262	.263
μ_4	.324	.378	.751	.810	.810	.796

regressors $u_{i,1}, \ldots, u_{i,q}$. Which subset of these ought to be included, and which ways are there of averaging over all models? The natural framework is that of $Y_i = \alpha + x_i^t\beta + u_i^t\gamma + \varepsilon_i$ for $i = 1, \ldots, n$, where the ε_i's are independent and $N(0, \sigma^2)$. Suppose that the u_i's have been made orthogonal to the x_i's, in the sense that $n^{-1}\sum_{i=1}^{n} x_i u_i^t = 0$. Then

$$J_{n,\text{full}} = \sigma^{-2}\text{diag}(2, \Sigma_{00}, \Sigma_{11}),$$

$$\text{with} \quad J_{n,\text{full}}^{-1} = \sigma^{-2}\text{diag}\left(\tfrac{1}{2}, \Sigma_{00}^{-1}, \Sigma_{11}^{-1}\right),$$

where $\Sigma_{00} = n^{-1}\sum_{i=1}^{n} x_i x_i^t$ and $\Sigma_{11} = n^{-1}\sum_{i=1}^{n} u_i u_i^t$. Inside this framework we may now study model selection and model averaging for different focus parameters, using methods developed in earlier sections. For the arguably most important case of $\mu = E(Y \mid x, u)$ at some given location (x, u), FMA estimators take the form

$$\hat{\mu}(x, u) = \sum_S c(S \mid D_n)(x^t\hat{\beta}_S + u_S^t\hat{\gamma}_S) = x^t\beta^* + u^t\gamma^*,$$

where the β_j^*'s and γ_k^*'s involved are nonlinear regression coefficient estimates. Methods and results of earlier sections can be used to settle on weighting schemes here, along with proper analysis of performances.

7. RISK COMPARISON

We are now in position to compare various model selection estimation and model averaging methods in terms of performance.

7.1 Comparing Risks in the Limit Experiment

In a given situation the risk function $nE(\hat{\mu} - \mu_{\text{true}})^2$ can be a quite complicated quantity, particularly when the estimator in question uses nonlinear weight schemes and when the underlying models are difficult. There is a drastic reduction in complexity as n grows, however, as spelled out in Section 4, in that the limiting risk $\tau_0^2 + R(\delta)$ depends on only a few crucial quantities. This allows broad comparisons to be made in a fairly easy fashion, by computing risk functions for situations and estimation schemes of interest, in their reduced limit experiment form.

It is often convenient to discuss performance in terms of $\overline{R}(a)$ instead of $R(\delta)$, because $a = K^{-1/2}\delta$ is scale independent with $Z \sim N_q(a, I)$. Note in this connection Remark 5.1 about reparameterization, which makes it possible to have K diagonal if one works with submodels represented by subsets of $(a_1, \ldots, a_q)^t$. Also note that, when K is diagonal,

$$\overline{R}(a) = \sum_{i,j} \omega_i\omega_j k_i^{1/2} k_j^{1/2}\left[V_{i,j}(a) + \{M_i(a) - a_i\}\{M_j(a) - a_j\}\right]$$
$$(7.1)$$

in terms of the means $M_i(a)$ of $W_i(Z)Z_i$ and covariance $V_{i,j}(a)$ of $W_i(Z)Z_i$ with $W_j(Z)Z_j$. This follows from Remark 4.3 and shows that even complicated risk functions may be computed easily via simulation. Before we go on to a briefly annotated list of estimators, we mention one more fact, namely, that for the simplest case of $q = 1$ model extension,

$$\overline{R}(a) = K\omega^2 R^*(a) \qquad \text{in terms of}$$

$$R^*(a) = E\{W(Z)Z - a\}^2. \quad (7.2)$$

risk R(a) for six methods

Figure 4. Risk Functions R(a) Associated With Six FMA Methods for q = 1 as in (7.2). These are symmetric around 0 and are displayed here for a ∈ [0,5]. The AIC (dotted line) starts at .572 with max-risk 1.650. The pretest approach, which uses .05 as a test level (dotted line), starts at .279 with high max-risk 2.464. The smoothed AIC (solid line) starts at .378 with max-risk 1.551. The empirical Bayes singleton method (dashed line) starts at .333 with max-risk 1.491. The method corresponding to (5.5) (solid line) starts at .467 with low max-risk 1.252. Finally, the method of Section 5.6 (dashed line), with $p_0 = .25$, starts at .335 with max-risk 1.619.

This is the one-dimensional risk function for the estimator $W(Z)Z$ for a in the standard experiment where $Z \sim N(a, 1)$. Here $1 - W(Z_n)$ and $W(Z_n)$ are the weights given to $\hat{\mu}_{\text{narr}}$ and $\hat{\mu}_{\text{full}}$. Such $R^*(a)$ functions are displayed in Figure 4 for various competing schemes.

1. *Narrow estimation.* Here $\bar{c}(\varnothing \mid Z) = 1$ and $\bar{c}(S \mid Z) = 0$ for other subsets, reflecting the optimistic belief, or blissful ignorance, that $a = 0$. The limiting risk is $\overline{R}_{\text{narr}}(a) = (\omega^t\delta)^2 = (\omega^t K^{1/2}a)^2$, which is unbounded and quickly becomes big in size. The risk is satisfactorily small when $\|a\|$ is small (and in that case for all estimands μ) or in cases where a is nearly orthogonal to $K^{1/2}\omega$ (which depends on the estimand).

2. *Wide model estimation.* Here $\bar{c}(S \mid Z) = 1$ for the full set, which leads to a constant minimax risk, $\overline{R}_{\text{full}}(a) = \omega^t K\omega$. This is satisfactory performance in some situations, but the estimator is often too guardedly pessimistic, losing out to methods that take into account that a could be small in size or have low correlation with $K^{1/2}\omega$ (making in that case $\psi = \omega^t\delta$ small in size).

3. *Hard and smooth AIC selection estimators.* The following comments are valid for the nonnested case. For simplicity of illustration also take K to be diagonal, in which case the AIC scores can be written as $\sum_{j \in S}(Z_j^2 - 2)$ in terms of $Z_j = D_j/k_j^{1/2}$. This entails a quite simple structure for the R_S regions, as the winning S is $\{j : |z_j| > \sqrt{2}\}$. Turning this around, one sees that

$$R_S = \left\{z : |z_j| > \sqrt{2} \text{ for each } j \in S \text{ and } |z_j| \leq \sqrt{2} \text{ for each } j \notin S\right\}.$$

In particular, $R_\varnothing = [-\sqrt{2}, \sqrt{2}]^q$ is the set inside which the method selects the narrow model. This also leads to $W_j(z) = I\{|z_j| > \sqrt{2}\}$ in (4.5), making it possible to calculate $\overline{R}(a)$

of (7.1) explicitly. We have done this in some further numerical comparison work, not reported on here due to limitations of space. The smoothed AIC scheme described in Section 5.2 uses weights proportional to $\exp\{\frac{1}{2}\sum_{j\in S}(z_j^2 - 2)\}$, as opposed to the hard thresholding $I\{|z_j| > \sqrt{2}\}$ involved in ordinary AIC.

4. *Hard and smooth FIC selection estimators.* The limiting risk functions for the two compromise estimators that use respectively the AIC and the FIC become $\tau_0^2 + \overline{R}_{\text{aic}}(a)$ and $\tau_0^2 + \overline{R}_{\text{fic}}(a)$, where $\overline{R}_{\text{fic}}(a) = \text{E}\{\omega^{\text{t}}K^{1/2}Z I_{\text{fic}}(Z) - \omega^{\text{t}}K^{1/2}a\}^2$ with an analogous definition for $\overline{R}_{\text{aic}}(a)$. Here $I_{\text{fic}}(z) = I\{(\omega^{\text{t}}K^{1/2}z)^2 \geq 2\omega^{\text{t}}K\omega\}$ and $I_{\text{aic}}(z) = I\{z^{\text{t}}z \geq 2q\}$. Investigations reported on in Claeskens and Hjort (2003) show that the FIC method often does better than the AIC. Also, it typically pays off to smooth the FIC weights as in (5.4).

5. *Average-across-singleton estimator.* When K is diagonal the method developed in Section 5.5 has

$$\overline{R}(a) = \text{E}\left[\sum_{j=1}^{q} \omega_j k_j^{1/2}\left\{\frac{\exp(\frac{1}{2}\hat{\rho}Z_j^2)}{\sum_{i=1}^{q}\exp(\frac{1}{2}\hat{\rho}Z_i^2)}\hat{\rho}Z_j - a_j\right\}\right]^2,$$

where $\hat{\rho} = \hat{\tau}^2/(1+\hat{\tau}^2)$ and $\hat{\tau} = (\|Z\|^2 - q)_+^{1/2}$.

6. *Other empirical Bayes schemes.* The main method of Section 5.5 may be analyzed via the appropriate

$$W_j(z) = \hat{\rho}\frac{p_1(1+\hat{\sigma}^2)^{-1/2}\exp(\frac{1}{2}\hat{\rho}z_j^2)}{p_0 + p_1(1+\hat{\sigma}^2)^{-1/2}\exp(\frac{1}{2}\hat{\rho}z_j^2)},$$

$$\text{with} \quad \hat{\rho} = \frac{\hat{\sigma}^2}{1+\hat{\sigma}^2},$$

along with simulation-based computation of $\overline{R}(a)$ as per (7.1). Alternatives may be analyzed similarly.

We have studied risk functions for various procedures for the one- and two-dimensional cases, but cannot report in any depth here due to limitations of space. Some brief remarks are as follows. (a) It pays to smooth the hard-core AIC and FIC methods, as in Sections 5.2 and 5.3, and sometimes with a κ bigger than 1 in (5.4). These risk functions are smaller than the constant minimax risk $\omega^{\text{t}}K\omega$ in a decent neighborhood around 0, then increase, and level off toward the minimax value as $\|a\|$ grows. (b) The singleton method of (5.5) does quite well in a reasonable neighborhood around 0 and along axes, where one $|\delta_j|$ is big but the others small, but its risk may become large when more than one $|\delta_j|$ becomes big. (c) The empirical Bayes scheme of Section 5.6, along with similarly inspired versions, does quite well in terms of low max-risk and being smaller than $\omega^{\text{t}}K\omega$ in a broad neighborhood around 0.

7.2 Risk Comparison in a Simulated Poisson Setting

Here we illustrate the mean squared error (mse) behavior of model-averaged and post-model selection estimators for Poisson regression. The situation we study has the narrow model containing an intercept only, whereas in the widest model four variables are included. In other words, counts Y_i are independent and Poisson with parameters ξ_i, where $\xi(u_i) = \exp(\beta_0 + \sum_{j=1}^{4}\gamma_j u_{i,j})$. There are at the outset $2^q = 16$ different submodels to consider, corresponding to inclusion or not of the four γ_j's. In the simulation study we took $\beta_0 = 1$

Table 4. Poisson Regression With $q = 4$ Extra Variables

Estimator	Setting (a)			Setting (b)		
	$n = 50$	$n = 200$	Limit	$n = 50$	$n = 200$	Limit
Post-AIC	20.26	16.65	17.85	12.13	17.74	14.22
Smooth AIC	14.67	13.24	13.73	9.89	13.49	11.27
Post-FIC	11.94	10.48	10.86	9.88	11.27	9.87
Smooth FIC	8.53	8.19	8.20	7.41	8.22	7.62
Wide model	15.47	12.11	12.56	10.01	12.03	10.11
Testing	16.57	14.97	20.17	11.85	16.06	16.55

NOTE: Simulated $n$mse values for setting (a), $u = (1, -.9, -.9, 1)^{\text{t}}$, and (b), $u = (1, 1, -.6, -.6)^{\text{t}}$.

and $\delta = (1, 1, 1, 1)^{\text{t}}$, that is, $\gamma_j = 1/\sqrt{n}$, and chose two focus points in the covariate space, $u = (1, -.9, -.9, 1)^{\text{t}}$ and $u = (1, 1, -.6, -.6)^{\text{t}}$. The four covariates $u_{i,1}, \ldots, u_{i,4}$ were taken to be independent and standard normal. Estimators were then simulated 1,000 times for each setting using the empirical $\widehat{J}$ matrix $n^{-1}\sum_{i=1}^{n}\hat{\xi}(u_i)u_i u_i^{\text{t}}$, with its submatrices and corresponding K matrix, and the ω vector calculated for each choice of u as $\hat{\omega} = \hat{\xi}(u)(\widehat{J}_{10}/\widehat{J}_{00} - u)$. Although the main point here is to compare methods for finite sample sizes, we also include in the table the population quantities $\tau_0^2 + R(\delta)$, which by Theorem 4.1 are the limits of n times mse. To compute these, we use the fact that $J_{\text{full}} = \exp(\beta_0)I_5$, from properties of the normal covariate distribution, in terms of the 5×5 identity matrix, and which entails $K = \exp(-\beta_0)I_4$. We then evaluated $R(\delta)$ by taking the average of a full million simulated versions of $\{\omega^{\text{t}}\hat{\delta}(D) - \omega^{\text{t}}\delta\}^2$, with $\hat{\delta}(D)$ as in (4.2) and $D \sim \text{N}_4(\delta, K)$.

Table 4 shows simulated $n$mse values for the following estimators: post-model selection using AIC and FIC; model-averaged estimators using smoothed AIC and FIC weights as in (5.2) and (5.4), the latter with $\kappa = 1$; the wide model estimator; and a testing approach where each variable is tested individually at a 5% level and only the significant variables are kept in the final model. Two sample sizes are used, $n = 50$ and $n = 200$. From the table it is observed that the model-averaged estimators have much lower mse values than the corresponding post-model selection estimators, which select one single model. For these settings, the FIC yields significantly smaller mse values than the AIC. Model averaging also performs better in terms of mse than the wide model method and outperforms the simple testing approach. For sample size 200, the simulated values are already close to the simulated theoretical mse values, confirming the theoretical derivations.

Table 5 shows for the same settings simulated coverage probabilities for $\xi(u)$ at a nominal 95% level for sample sizes n

Table 5. Poisson Regression With $q = 4$ Extra Variables

Estimator	Setting (a)			Setting (b)		
	$n = 50$	$n = 100$	$n = 200$	$n = 50$	$n = 100$	$n = 200$
Post-AIC	.935	.946	.948	.941	.948	.946
Smooth AIC	.935	.946	.948	.942	.946	.946
Post-FIC	.933	.946	.955	.939	.946	.949
Smooth FIC	.957	.967	.974	.957	.968	.971
Buckland et al.	.925	.926	.929	.928	.897	.916
Wide model	.936	.947	.948	.944	.947	.946
Testing	.815	.833	.828	.846	.768	.786
Naive AIC	.773	.827	.820	.833	.732	.783
Naive BIC	.720	.700	.690	.748	.568	.638

NOTE: Simulated coverage probabilities for $\xi(u)$ for setting (a), $u = (1, -.9, -.9, 1)^{\text{t}}$, and (b), $u = (1, 1, -.6, -.6)^{\text{t}}$.

equal to 50, 100, and 200, based on 10,000 simulation replicates of all estimators. The table includes first results of application of definitions (4.8) with $D_n = \sqrt{n}\hat{\gamma}_{\text{full}}$ for the first four methods shown: post-model selection using AIC and FIC (with $\kappa = 1$) and their smoothed model-averaged versions. The approach by Buckland et al. (1997) used the same smoothed AIC weights, but then used the standard error estimate described in Section 4.3 in conjunction with a simple nonbiased normal approximation. The table furthermore displays results for the confidence interval coming from using the widest model estimator and the testing strategy explained previously, both employing standard normal percentiles. It also shows the results of what happens to the coverage probability when ignoring the model selection step after AIC or BIC model selection.

With increasing sample size the corrected versions approach the nominal level of 95%, although the smoothed FIC values are a little larger than the nominal value, at least for this setting. By construction, the wide method is the safest method and will produce asymptotically correct confidence intervals. As a consequence of using an imperfect distributional approximation, the method by Buckland et al. does not reach nominal coverage in the performed simulations. The testing approach using normal percentiles produces confidence intervals with significantly lower than nominal coverage values. And as expected from theoretical considerations, see Section 4.3, ignoring model selection results in confidence intervals with too low coverage probabilities, as is illustrated by the last two rows in the table.

8. GENERALIZED RIDGING: SHRINKING IN PARAMETRIC MODELS

The development of Section 4 gave an instructive bridge from compromise estimators $\hat{\mu}$ of type (4.1) to estimators of $\psi = \omega^t\delta$ of type (4.6). For some purposes the class of (4.1) estimators is not quite large enough, however. For example, Theorem 4.1 does not cover the full class of natural $\hat{a}_j(z) = \omega_j k_j^{1/2} W_j(Z) Z_j$-type estimators encountered as a consequence of exploiting this theorem; see Remark 4.3. This section expands the horizon by proposing and investigating certain generalized ridge estimators, which shrink the $\hat{\gamma}_S$ estimators toward γ_0. Such shrinking may be particularly beneficial when the number q of extra parameters is moderate or growing compared to a fixed number p of core parameters θ. A quite general class of BMA estimators will, in fact, behave just in this way, as shown in Section 9.

The intention is to stick to the narrow model as a form of basis, but to consider downweighting aspects of the more risky γ extensions, via estimators of the form $\tilde{\gamma}_S$ that shrink the $\hat{\gamma}_S$ toward γ_0, with an amount somehow dictated by D_n. The idea is to use $(\hat{\theta}_S, \tilde{\gamma}_S, \gamma_{0,S^c})$ as estimators in the S subset model, and, more specifically,

$$\tilde{\mu}_S = \mu(\hat{\theta}_S, \tilde{\gamma}_S, \gamma_{0,S^c}),$$

$$\text{where} \quad \tilde{\gamma}_S - \gamma_{0,S} = \{1 - \varepsilon_S(D_n)\}(\hat{\gamma}_S - \gamma_{0,S^c})$$

for suitable functions $\varepsilon_S(d)$. The cases encountered earlier correspond to these functions being identically 0. For these estimators

$$\begin{pmatrix} \sqrt{n}(\hat{\theta}_S - \theta_0) \\ \sqrt{n}(\tilde{\gamma}_S - \gamma_{0,S}) \end{pmatrix} \xrightarrow{d} \begin{pmatrix} C_S \\ \{1 - \varepsilon_S(D)\}D_S \end{pmatrix},$$

which leads to

$$\sqrt{n}(\tilde{\mu}_S - \mu_{\text{true}})$$

$$\xrightarrow{d} \Lambda_S = \left(\frac{\partial\mu}{\partial\theta}\right)^t C_S + \left(\frac{\partial\mu}{\partial\gamma_S}\right)^t \{1 - \varepsilon_S(D)\}D_S - \left(\frac{\partial\mu}{\partial\gamma}\right)^t \delta.$$

With some algebraic work we find

$$\Lambda_S = \left(\frac{\partial\mu}{\partial\theta}\right)^t J_{00}^{-1} M + \omega^t\{(I - K^{-1/2}H_S K^{1/2})^t\delta$$

$$- K^{1/2}H_S K^{-1/2}W\}$$

$$- \varepsilon_S(D)\left(\frac{\partial\mu}{\partial\gamma}\right)^t K^{1/2}H_S K^{-1/2}(\delta + W).$$

To work with estimator-after-selection estimators, we are again led to consider the class of estimators

$$\tilde{\mu} = \sum_S c(S \mid D_n)\tilde{\mu}_S = \sum_S c(S \mid D_n)\mu(\hat{\theta}_S, \tilde{\mu}_S, \gamma_{0,S^c}), \quad (8.1)$$

with coefficients summing to 1 and allowed to depend on D_n of (3.1). There is a limiting distribution for $\sqrt{n}(\tilde{\mu} - \mu_{\text{true}})$ admitting the representation

$$\Lambda = \left(\frac{\partial\mu}{\partial\theta}\right)^t J_{00}^{-1} M + \omega^t\{\delta - G(D)^t(\delta + W)\}$$

$$- \left(\frac{\partial\mu}{\partial\gamma}\right)^t G^*(D)^t(\delta + W),$$

where $G(D) = \sum_S c(S \mid D)K^{-1/2}H_S K^{1/2}$ and $G^*(D) = \sum_S c(S \mid D)\varepsilon_S(D)K^{-1/2}H_S K^{1/2}$. A special case of interest is when each $\varepsilon_S(D)$ the same, corresponding to equal ridging in all γ_j directions. Then $G^*(D) = \varepsilon(D)G(D)$.

These efforts lead to a limiting risk expression for estimators of the form (8.1), namely,

$$E\Lambda^2 = \tau_0^2 + E\left[\omega^t\{\delta - G(D)^t D\} - \left(\frac{\partial\mu}{\partial\gamma}\right)^t G^*(D)^t D\right]^2$$

$$= \tau_0^2 + E(\tilde{\psi} - \omega^t\delta)^2,$$

say, where the estimator of $\psi = \omega^t\delta$ this time becomes

$$\left\{\omega^t G(D)^t + \left(\frac{\partial\mu}{\partial\gamma}\right)^t G^*(D)^t\right\}D$$

$$= \left[\left(\frac{\partial\mu}{\partial\theta}\right)^t J_{00}^{-1} J_{01} G(D)^t - \left(\frac{\partial\mu}{\partial\gamma}\right)^t \{G(D) - G^*(D)\}^t\right]D.$$

This would take the form

$$\tilde{\psi} = \left[\left(\frac{\partial\mu}{\partial\theta}\right)^t J_{00}^{-1} J_{01} - \{1 - \varepsilon(D)\}\left(\frac{\partial\mu}{\partial\gamma}\right)^t\right]G(D)^t D$$

in the case of the same $\varepsilon(D)$ for all subsets under consideration. We see that this bridges from the situation of Theorem 4.1, for $\varepsilon(D) = 0$, to the result for the narrow procedure, for $\varepsilon(D) = 1$. It also shows that using $\varepsilon(D)$ may shrink the size of the $\partial\mu/\partial\gamma$ component here, in its turn often lowering the variance level.

9. A FREQUENTIST VIEW OF BMA

Bayesian model averaging essentially amounts to putting down prior probabilities $p(S)$ for all submodels, prior distributions $\pi_S(\theta, \delta_S)$ for the parameters inside the S submodel, and then applying Bayes's theorem suitably. The basics of such machineries is covered in Draper (1995) and Hoeting et al. (1999). In our framework the posterior of the parameters may be expressed as

$$\pi_n(\theta, \delta) = \sum_S p_n(S)\pi_{n,S}(\theta, \delta_S). \qquad (9.1)$$

Here $\pi_{n,S}(\theta, \delta_S)$ is the posterior calculated under the S model (in particular, then $\delta_j = 0$ for $j \notin S$), whereas $p_n(S) = p(S)\lambda_n(S)/\sum_{S'} p(S')\lambda_n(S')$ is the probability of model S, given the data. Here

$$\lambda_n(S) = \int L_{n,S}(\theta, \gamma_0 + \delta_S/\sqrt{n})\pi_S(\theta, \delta_S)\, d\theta\, d\delta_S \qquad (9.2)$$

is the integrated likelihood of model S, involving the likelihood $L_{n,S}$ for this model. The $\lambda_n(S)$ is also the marginal distribution at the observed data. We will derive approximations and precise limit distribution results for the quantities involved in (9.1) and (9.2), under our local alternative framework. The limit behavior of BMA schemes has apparently not been studied before.

9.1 The Posterior Model Probabilities

We need to understand the behavior of $\lambda_n(S)$. The familiar BIC statistic stems, in fact, from an approximation to this quantity. To review and comment on this approximation, let, as before, $\hat{\theta}_S$ and $\hat{\delta}_S = \sqrt{n}(\hat{\gamma}_S - \gamma_{0,S})$ be the maximum likelihood estimators inside the S model. Then

$$\lambda_n(S) \doteq L_{n,S}(\hat{\theta}_S, \hat{\gamma}_S)n^{-(p+|S|)/2}$$
$$\times (2\pi)^{(p+|S|)/2}|J_{n,S}|^{-1/2}\pi_S(\hat{\theta}_S, \hat{\delta}_S) \quad (9.3)$$

is one possible approximation. Here $J_{n,S} = -n^{-1} \times \sum_{i=1}^n \partial^2 \log f(Y_i, \hat{\theta}_S, \hat{\gamma}_S)/\partial\alpha_S \partial\alpha_S^t$ is the observed information matrix of size $(p + |S|) \times (p + |S|)$, using α_S to denote the parameter vector with θ and γ_S. The consequent $2\log\lambda_n(S) \approx 2\max\log L_{n,S} - (p + |S|)\log n$ is often called "the BIC approximation"; see, for example, Hoeting et al. [1999, eq. (13), modulo an incorrect constant]. Claim (9.3) may be proved using arguments similar to those needed to show Proposition 9.1.

It is important to note, however, that the asymptotic approximation (9.3), which underlies the BIC, is valid in the framework of fixed models $f(y, \theta, \gamma)$ and a fixed $f_{\text{true}}(y)$, and where, in particular, also $\delta = \sqrt{n}(\gamma - \gamma_0)$ grows with n. In such a framework the best model will win in the end; that is, the candidate model S_0 with smallest Kullback–Leibler distance to the true density will have $p_n(S_0) \to 1$ as n grows. This follows as the dominant term of $\max\log L_{n,S}$ will be n times $\max \int f_{\text{true}}(y)\log f(y, \theta, \gamma)\, dy$. In our framework of local alternative models, the magnifying glass is focused on the $\sqrt{n}(\gamma - \gamma_0)$ scale, and different results apply. Maximized log-likelihoods are then not $O_p(n)$ apart, as under the fixed model scenario, but have differences related to noncentral chi-squared distributions. Second, the $n^{-|S|/2}$ ingredient, crucial to the BIC, disappears.

For the following result, which provides a more accurate approximation than the BIC-related (9.3) when the scale of model departures from the narrow model is that of $\delta = \sqrt{n}(\gamma - \gamma_0)$, we let $\phi(\cdot, \Sigma)$ denote the density of a $N(0, \Sigma)$.

Proposition 9.1. Let the prior for the S subset model take the form $\pi_0(\theta)\pi_S(\delta_S)$, with π_0 continuous in a neighborhood around θ_0. Then, under standard regularity conditions, when n grows,

$$\lambda_n(S) \doteq L_{n,S}(\hat{\theta}_S, \hat{\gamma}_S)n^{-p/2}$$
$$\times (2\pi)^{(p+|S|)/2}\pi_0(\hat{\theta}_S)|J_{n,S}|^{-1/2}\kappa_n(S),$$

where $\kappa_n(S) = \int \phi(\delta_S - \hat{\delta}_S, J_{n,S}^{11})\pi_S(\delta_S)\, d\delta_S$. The approximation holds in the sense that $\log\lambda_n(S)$ is equal to the logarithm of the right side plus a remainder term of size $O_p(n^{-1/2})$. Also, $J_{n,S}^{11}$ is the lower right-hand $|S| \times |S|$ submatrix of $J_{n,S}^{-1}$.

When n grows we also have $J_{n,S} \to_p J_S$, defined in Section 3.1, and the limit of $J_{n,S}^{11}$ is $K_S = (\pi_S K^{-1}\pi_S^t)^{-1}$. Combining this with some previous results, reached in conjunction with (3.4), we find

$$\lambda_n(S) \doteq \text{const.} \exp\left(\tfrac{1}{2}\hat{\delta}_S^t K_S^{-1}\hat{\delta}_S\right)(2\pi)^{|S|/2}|J_S|^{-1/2}$$
$$\times \int \phi(\delta_S - \hat{\delta}_S, K_S)\pi_S(\delta_S)\, d\delta_S,$$

where the constant in question is $n^{-p/2}(2\pi)^{p/2}\pi_0(\hat{\theta})$. This also leads to a precise description of posterior probabilities for the different models in the canonical limit experiment. This is the situation of large n where all quantities have been estimated with full precision except δ, for which we must be content with the limit $D \sim N(\delta, K)$ of $D_n = \sqrt{n}(\hat{\gamma}_{\text{full}} - \gamma_0)$. Here $p(S \mid D) \propto p(S)\lambda(S)$, where

$$\lambda(S) = \exp\left(\tfrac{1}{2}D_S^t K_S^{-1}D_S\right)(2\pi)^{|S|/2}|J_S|^{-1/2}$$
$$\times \int \phi(\delta_S - D_S, K_S)\pi_S(\delta_S)\, d\delta_S$$
$$= \exp\left(\tfrac{1}{2}\text{AIC}_S\right)\exp(|S|)(2\pi)^{|S|/2}|J_S|^{-1/2}$$
$$\times \int \phi(\delta_S - D_S, K_S)\pi_S(\delta_S)\, d\delta_S$$

and $D_S = K_S\pi_S K^{-1}D$. We use here $\text{AIC}_S = D_S^t K_S^{-1}D_S - 2|S|$ from Section 3.3.

9.2 Bayesian Model Choice With the Canonical Normal Priors

The primary special case is when δ_S has the prior $N(0, \tau_S^2 K_S)$. This corresponds to independent and equally spread-out priors around 0 for the transformed parameters $a_S = \pi_S a$, where $a = K^{-1/2}\delta$ on the canonical scale, and where again $Z \sim N_q(a, I)$. Then

$$\lambda(S) = \exp\left(\frac{1}{2}\frac{\tau_S^2}{1+\tau_S^2}D_S^t K_S^{-1}D_S\right)$$
$$\times (1+\tau_S^2)^{-|S|/2}|J_{00}|^{-1/2}. \quad (9.4)$$

The last determinant is independent of $|S|$ and emerges via $|J_S|^{-1/2}|K_S|^{-1/2}$, because $|J_S| = |J_{00}| |K_S|^{-1}$. Result (9.4) is

also valid for $S = \varnothing$, corresponding to the narrow model, for which $\lambda(\varnothing) = |J_{00}|^{-1/2}$.

This also gives rise to a new Bayesian information criterion, which we may term the BLIC, with L for "local," reminding us of the local model extension framework (2.2). This criterion is reached by following the original BIC path, but using a different statistical magnifying glass, focusing on $\gamma_0 + \delta/\sqrt{n}$-type neighboring models. From (9.4) our criterion reads

$$\text{BLIC} = \frac{\tau_S^2}{1 + \tau_S^2} D_S^t K_S^{-1} D_S - |S| \log(1 + \tau_S^2) + 2\log p(S),$$

because the posterior model probability is close to being proportional to $p(S)\lambda(S)$. Here τ_S is meant to be a spread measure for δ_S in submodel S and for the narrow model BLIC $= 2\log p(\varnothing)$. The candidate model with largest BLIC is the most probable one, given the data, in the Bayesian formulation and is selected.

The preceding formula is valid for the limit experiment. For real data we use $\hat{\delta}_S$ for D_S, leading to

$$\widehat{\text{BLIC}} = \frac{\tau_S^2}{1 + \tau_S^2} n(\hat{\gamma}_S - \gamma_{0,S})^t \widehat{K}_S^{-1}(\hat{\gamma}_S - \gamma_{0,S})$$
$$- |S| \log(1 + \tau_S^2) + 2\log p(S).$$

Furthermore, we may estimate the spread. First, $D_S^t K_S^{-1} D_S$, given δ, is a noncentral chi-squared with parameter $\delta_S^t K_S^{-1} \delta_S$. Taking the mean of $|S| + \delta_S^t K_S^{-1} \delta_S$ again gives $|S|(1 + \tau_S^2)$. We may thus suggest $1 + \tau_S^2$ estimated, in this empirical Bayes fashion, by $D_S^t K_S^{-1} D_S/|S|$. This gives say

$$\text{BLIC}^* = |S|\{\hat{\tau}_S^2 - \log(1 + \hat{\tau}_S^2)\} + 2\log p(S),$$
$$\text{with} \quad \hat{\tau}_S^2 = \max\{D_S^t K_S^{-1} D_S/|S| - 1, 0\}.$$

Various alternatives may also be considered.

9.3 Posteriors in Submodels

We need to investigate the behavior of the posterior distributions $\pi_{n,S}(\theta, \delta_S)$, conditional on model S, and, in particular, their means $\tilde{\mu}_S = \text{E}_S(\mu \mid \text{data})$. It will become clear that, for large n, the distribution of θ will be tightly concentrated around $\hat{\theta}_S$, whereas the part of the prior related to δ_S will not be "washed away" by the data. This is because the chimeric parameter δ will not be consistently estimated as data accumulate; the best we may do is via $\hat{\delta}_{\text{full}} \xrightarrow{d} N_q(\delta, K)$. The posterior for δ_S will, in fact, go to

$$\pi_S(\delta_S \mid \hat{\delta}_S) = \text{const.} \, \pi_S(\delta_S) \exp\{-\tfrac{1}{2}(\delta_S - \hat{\delta}_S)^t K_S^{-1}(\delta_S - \hat{\delta}_S)\}$$
$$= \text{const.} \, \pi_S(\delta_S)\phi(\delta_S - \hat{\delta}_S, K_S), \quad (9.5)$$

as shall be seen later. Thus, $\text{E}(\delta_S \mid \text{data})$ is for large n essentially a function of $\hat{\delta}_S$, which again is a function of Z_n of (4.4), per Section 3.3. In the limit experiment, where $\hat{\delta}_S \xrightarrow{d} D_S = K_S \pi_S K^{-1}(\delta + W)$ with mean $K_S \pi_S K^{-1}\delta$ and variance matrix K_S, write

$$\text{E}_L(\delta_S \mid D_S)$$
$$= \frac{\int \delta_S \pi_S(\delta_S) \exp\{-\tfrac{1}{2}(\delta_S - D_S)^t K_S^{-1}(\delta_S - D_S)\} \, d\delta_S}{\int \pi_S(\delta_S) \exp\{-\tfrac{1}{2}(\delta_S - D_S)^t K_S^{-1}(\delta_S - D_S)\} \, d\delta_S}.$$

We then have the following extension of Lemma 3.3.

Proposition 9.2. Under the conditions of the previous proposition, the Bayesian submodel estimator $\tilde{\mu}_S = \text{E}_S(\mu \mid \text{data})$ is asymptotically equivalent to the simpler estimator $\bar{\mu}_S = \text{E}\{\mu(\hat{\theta}_S, \gamma_{0,S} + \delta_S/\sqrt{n}) \mid \hat{\delta}_S\}$, where the distribution in question is that of (9.5). Also,

$$\sqrt{n}(\tilde{\mu}_S - \mu_{\text{true}})$$
$$\xrightarrow{d} \tilde{\Lambda}_S = \left(\frac{\partial \mu}{\partial \theta}\right)^t C_S + \left(\frac{\partial \mu}{\partial \gamma_S}\right)^t \text{E}_L(\delta_S \mid D_S) - \left(\frac{\partial \mu}{\partial \gamma}\right)^t \delta.$$

9.4 BMA Approximations

The approximations to $\tilde{\mu}_S$ indirectly touched on here are of separate value. The simplest of these, from the second half of the proof, is $\hat{\mu}_S - (\frac{\partial \mu}{\partial \gamma_S})^t\{\hat{\delta}_S - \text{E}(\delta_S \mid \hat{\delta}_S)\}$. It is also useful to record an approximation to the conditional variance $\tilde{\sigma}_S^2 = \text{Var}_S(\mu \mid \text{data})$. One first may show that $\tilde{\sigma}_S^2 = \text{E}\{\mu(\hat{\theta}_S, \gamma_{0,S} + \delta_S/\sqrt{n})^2 \mid \hat{\delta}_S\} - \tilde{\mu}^2 + o(n^{-1})$, and is then via (9.5) and renewed Taylor expansion led to $\tilde{\sigma}_S^2 \doteq n^{-1}(\frac{\partial \mu}{\partial \gamma_S})^t \text{Var}(\delta_S \mid \hat{\delta}_S)\frac{\partial \mu}{\partial \gamma_S}$.

The primary special case here is again the normal priors for δ_S studied in Section 8.2. Then the posterior is normal with mean $\rho_S \hat{\delta}_S$ and variance $\rho_S K_S$, where $\rho_S = \tau_S^2/(1 + \tau_S^2)$. Thus,

$$\tilde{\mu}_S \doteq \hat{\mu}_S - \left(\frac{\partial \mu}{\partial \gamma_S}\right)^t \hat{\delta}_S(1 - \rho_S),$$
$$\tilde{\sigma}_S^2 \doteq n^{-1}\left(\frac{\partial \mu}{\partial \gamma_S}\right)^t K_S \frac{\partial \mu}{\partial \gamma_S}\rho_S.$$

Also, from the proposition,

$$\tilde{\Lambda}_S = \left(\frac{\partial \mu}{\partial \theta}\right)^t C_S + \rho_S\left(\frac{\partial \mu}{\partial \gamma_S}\right)^t D_S - \left(\frac{\partial \mu}{\partial \gamma}\right)^t \delta.$$

But this is exactly as in Section 6, with shrinking factor $\varepsilon_S(D) = 1 - \rho_S = 1/(1 + \tau_S^2)$ independent of D. When τ_S is small, the prior is informative and tight, the shrinkage is high, and the Bayes estimator is in the $\tau_S \to 0$ limit the same as the narrow estimator $\hat{\mu}_{\text{narr}}$. If, on the other hand, τ_S becomes big, then the prior is diffuse and the shrinkage is small; in the limit case $\tau_S \to \infty$, the Bayes estimator is the same as the maximum likelihood estimator $\hat{\mu}_S$.

For BMA estimators the limiting risk function to study is $R(\delta) = \text{E}(\tilde{\psi} - \omega^t \delta)^2$, where

$$\tilde{\psi} = \omega^t G(D)^t D + \left(\frac{\partial \mu}{\partial \gamma}\right)^t G^*(D)^t D,$$

$G(D) = \sum_S c(S \mid D)K^{-1/2}H_S K^{1/2}$ and $G^*(D) = \sum_S c(S \mid D) \times (1 + \tau_S^2)^{-1}K^{-1/2}H_S K^{1/2}$. Furthermore, $c(S \mid D)$ is proportional to $p(S)\lambda(S)$, with $\lambda(S)$ as in (9.4).

A result analogous to (9.1) holds for the posterior distribution of $\mu = \mu(\theta, \gamma_0 + \delta/\sqrt{n})$, which we write as $\pi_n(\mu) = \sum_S p_n(S)\pi_{n,S}(\mu)$. The Bayes estimator (under quadratic loss) becomes $\tilde{\mu} = \text{E}(\mu \mid \text{data}) = \sum_S p_n(S)\tilde{\mu}_S$, whereas $\text{Var}(\mu \mid \text{data})$, the natural Bayesian measure of spread, becomes $\sum_S p_n(S)\{\tilde{\sigma}_S^2 + (\tilde{\mu}_S - \tilde{\mu})^2\}$. These formulas allows one to carry out approximate BMA analysis with simple computations, without, for example, MCMC computations.

10. CONCLUDING REMARKS

10.1 Amendments When the Largest Model Does Not Hold

Our machinery has been developed under the key assumption (2.2), which says that the true data generating mechanism should be inside the largest of the parametric models considered. Such an assumption may be tested via goodness-of-fit methods, but can never be established with certainty. Here we investigate briefly what happens when assumption (2.2) is not required to hold, relying on extended theory developed in Claeskens and Hjort (2003, sec. 8).

Assume that the true density for the data takes the form

$$f_{\text{true}}(y) = f(y, \theta_0, \gamma_0)\{1 + r(y)/\sqrt{n}\} + o(1/\sqrt{n}) \quad (10.1)$$

for a suitable $r(y)$ function, with $\int f_0|r|\,dy$ finite and $\int f_0 r\,dy = 0$, where $f_0(y) = f(y, \theta_0, \gamma_0)$. Condition (2.2) corresponds to the special case $r(y) = V(y)^t \delta$, with $V(y)$ as in Section 3.1. Because there are no "true parameters" now, consider instead the least false parameter, say $\mu_{\text{lf}} = \mu(\theta_n, \gamma_n)$. Here (θ_n, γ_n) are the least false parameters inside the $f(y, \theta, \gamma)$ family, that is, those minimizing what is the Kullback–Leibler distance $\int f_{\text{true}}(y) \log\{f_{\text{true}}(y)/f(y, \theta, \gamma)\}\,dy$. It is shown in Claeskens and Hjort (2003, sec. 8) that $\theta_n = \theta_0 + \eta_0/\sqrt{n}$ and $\gamma_n = \gamma_0 + \delta_0/\sqrt{n}$, apart from terms of smaller order, for constants η_0, δ_0 depending on $\int f_0 U r\,dy$ and $\int f_0 V r\,dy$, as explained there. It is also shown that

$$\sqrt{n}(\hat{\mu}_S - \mu_{\text{lf}}) \xrightarrow{d} \tilde{\Lambda}_S = \left(\frac{\partial \mu}{\partial \theta}\right)^t J_{00}^{-1} M$$
$$+ \omega^t\{\delta_0 - K^{1/2} H_S K^{-1/2}(\delta_0 + W)\}. \quad (10.2)$$

This is actually close to the result derived in Lemma 3.3, but now under wider start assumptions. The first point to note is that η_0 has dropped out; the second is that the agnostic parameter δ_0 takes the place of our earlier δ. Also, $D_n = \hat{\delta}_{\text{full}} \xrightarrow{d} D = \delta_0 + W \sim N_q(\delta_0, K)$, in generalization of (3.1). This means that the theory of Sections 4–6, about compromise, postselection, and shrinkage estimators, essentially goes through, with small amendments, and the methods developed are still in force. The difference is mostly related to interpretation, not to algorithms, so to speak; the precision of the estimators is interpreted and assessed in terms of the closeness of the agnostic μ_{lf}, rather than to the "true" focus parameter.

10.2 Breadth of Applications

It should be clear from our unified framework and application examples that there is a wide range of potential applications of our methods. Subset selection and model averaging can, in particular, be implemented and studied for quite general regression models, such as generalized linear models. Versions of our methods and results would also hold for models with dependence and for various stochastic process models. The essential requirement is that ordinary likelihood analysis should be valid, with limiting normality of the maximum likelihood estimators and so on. Our study indicates that it would be useful to carry out more extensive risk comparisons in the limit experiment, as touched on in Section 7.1, in that conclusions reached there will have implications for a fair range of situations.

10.3 Tolerance Radii

Sometimes ignorance is strength, and it may be better to stick to a simple model rather than going for a more complex one. This is captured well by our results of Sections 3 and 4. These may be used to characterize situations where a given S subset model gives better results than a competing S'. We find, in particular, that inference using the narrow model is better than that using the fullest model, provided $|\omega^t \delta| \leq (\omega^t K \omega)^{1/2}$. For a given estimand this describes a band of infinite length for δ. On the other hand, inside the ellipsoid where $\delta^t K^{-1} \delta \leq 1$ narrow model inference is better than full-model inference, for all estimands.

10.4 Two Uses of Models

Sometimes statistical modeling strives to come close to a superior scientific explanation of the phenomenon being studied, for example, in physics or biology. In this article we are employing models differently, as pragmatic approximations of reality, with the aim of generating estimates and predictions with good precision. See the engaging discussion of Breiman (2001) and section 1 of Claeskens and Hjort (2003) for further comments.

10.5 Optimal Methods

In addition to the methods proposed in Section 5, one may try to develop FMA schemes with suitable optimality properties. The Bayes methods we have discussed are optimal wrt the criterion of minimizing prior-weighted risk. The full-model estimator is the unique minimax estimator, under the $(\omega^t \delta - \omega^t \hat{\delta})^2$ loss function, with constant risk $\omega^t K \omega$. Other criteria might in one way or another involve ideas of restricting max-risk under the constraint of doing well at or near $\delta = 0$. Methods developed, for other purposes, in Bickel (1981, 1983, 1984) and Berger (1982) are of relevance here, but cannot be applied directly in that we restrict attention to FMA regimes.

10.6 Bootstrapping Does Not Work

To explain why bootstrapping cannot be relied upon in our model choice framework, consider the following situation. It is simple but representative of our general local model choice context. There are independent observations $Y_i \sim N(\mu, 1)$, where the narrow model holds that $\mu = 0$ and the wider model takes μ unknown; thus, $\hat{\mu}_{\text{full}} = \overline{Y}_n$. In the framework of local alternatives $\mu = \delta/\sqrt{n}$, where $Z_n = \sqrt{n}\overline{Y}_n$ is the natural test statistic, consider a model average estimator $\hat{\mu}$, which gives weight $1 - W(Z_n)$ to $\hat{\mu}_{\text{narr}}$ and weight $W(Z_n)$ to $\hat{\mu}_{\text{full}}$, that is, $\hat{\mu} = W(\sqrt{n}\overline{Y}_n)\overline{Y}_n$. First study

$$\Lambda_n = \sqrt{n}(\hat{\mu} - \mu_{\text{true}}) = W(\sqrt{n}\overline{Y}_n)\sqrt{n}\overline{Y}_n - \delta$$
$$\stackrel{d}{=} W(\delta + N)(\delta + N) - \delta,$$

where N represents a standard normal. Then study bootstrapped data Y_i^* from the estimated full model $N(\hat{\mu}_{\text{full}}, 1)$, with resulting bootstrap estimator $\hat{\mu}^* = W(\sqrt{n}\overline{Y}_n^*)\overline{Y}_n^*$. Here we find

$$\Lambda_n^* = \sqrt{n}(\hat{\mu}^* - \hat{\mu}) = W(\sqrt{n}\overline{Y}_n^*)\sqrt{n}\overline{Y}_n^* - \hat{\delta}$$
$$\stackrel{d}{=} W(\hat{\delta} + N')(\hat{\delta} + N') - \hat{\delta},$$

where N' represents another standard normal, independent of N given previously. Thus, the distributions of Λ_n and Λ_n^* are *not* close (excluding now the special case $W = 1$, which corresponds to using the wide estimator), because $\hat{\delta} = \sqrt{n}\hat{\mu}$ does not go to δ in probability.

10.7 Finite-Sample Correction and Non-ML Estimators

We have made extensive use of the first-order asymptotic theory for maximum likelihood estimation, yielding clear and concise descriptions of limit distributions and so on. Although this is already quite satisfactory, it is clear that suitable finite-sample corrections could be developed, perhaps for particular classes of models, in order to improve approximations. Work by Hurvich and Tsai (1989), extensively discussed in Burnham and Anderson (2002), is of relevance here. Another direction for future research is that of using robust estimators, perhaps of the M-estimator variety, instead of maximum likelihood estimators. It may also be important to use more robust weighting schemes.

APPENDIX: PROOFS OF LEMMAS AND THEOREMS

For the setup of Section 3.1, it is assumed that the log-density has two continuous partial derivatives around (θ_0, γ_0), so that

$$\log \frac{f(y, \theta_0 + s, \gamma_0 + t)}{f(y, \theta_0, \gamma_0)}$$

$$= \left(\begin{matrix} U(y) \\ V(y) \end{matrix}\right)^{\mathrm{t}} \left(\begin{matrix} s \\ t \end{matrix}\right) + \frac{1}{2}\left(\begin{matrix} s \\ t \end{matrix}\right)^{\mathrm{t}} W(y) \left(\begin{matrix} s \\ t \end{matrix}\right) + R(y, s, t) \quad \text{(A.1)}$$

for (s, t) small in $\mathcal{R}^{p+q}$, involving the matrix $W(y)$ of second log-density derivatives at the null point and a remainder term $R(y, s, t)$. It is also required that the variance matrix J_{full} of $(U(Y), V(Y))$ under $f_0(y) = f(y, \theta_0, \gamma_0)$, which is also the negative mean of $W(Y)$ under f_0, is finite and of full rank. This also gives rise to the representation $f(y, \theta_0, \gamma_0 + t) = f_0(y)\{1 + V(y)^{\mathrm{t}}t + R_2(y, t)\}$, where $R_2(y, t)$ is typically small enough to make $f_0(y)R_2(y, t)$ of order $o(\|t\|^2)$ uniformly in y, and to

$$f_{\text{true}}(y) = f_0(y)\{1 + V(y)^{\mathrm{t}}\delta/\sqrt{n} + R_2(y, \delta/\sqrt{n})\}. \quad \text{(A.2)}$$

Various sets of regularity conditions may now be put up to reach the desired conclusions, working either with (A.1) or (A.2) as convenient. Consider, in fact, the following assumptions.

(C1) The two integrals $\int f_0(y)U(y)R_2(y, t)\,dy$ and $\int f_0(y)V(y) \times R_2(y, t)\,dy$ are both $o(\|t\|)$.
(C2) The variables $|U_i^2 V_j|$ and $|V_i^2 V_j|$ have finite mean under f_0, for each i, j.
(C3) The two integrals $\int f_0(y)\|U(y)\|^2 R_2(y, t)\,dy$ and $\int f_0(y) \times \|V(y)\|^2 R_2(y, t)\,dy$ are both $o(1)$.
(C4) The log-density has three continuous derivatives wrt all $p + q$ parameters in a neighborhood around (θ_0, γ_0) and are there dominated by functions with finite means under f_0.

Conditions (C1) and (C3) are quite weak, and are implied by the stronger condition (C4). Then the integrals of (C1) and (C3) are in fact $o(\|t\|^2)$. Condition (C4) will hold for most models, as will (C2).

Proof of Lemma 3.1

Under conditions (C1), (C2), and (C3), this is accomplished via the multivariate Lindeberg theorem (which is the univariate Lindeberg theorem in conjunction with the Cramér–Wold device; see, e.g., Serfling 1980, sec. 1.9). Condition (C1) implies $EU(Y_i) = J_{01}\delta/\sqrt{n} + o(1/\sqrt{n})$ and $EV(Y_i) = J_{11}\delta/\sqrt{n} + o(1/\sqrt{n})$, whereas (C2) and (C3) ensure that the variance matrix of $\sqrt{n}(\overline{U}_n, \overline{V}_n)$ goes to J_{full}. The Lindeberg

requirement for asymptotic normality here demands integrals of $U_i^2 I\{\|U\| \geq \sqrt{n}\varepsilon\}$ and $V_i^2 I\{\|V\| \geq \sqrt{n}\varepsilon\}$ wrt the (A.2) density go to 0, and this is secured, for each positive ε, under (C2) and (C3).

Proof of Lemma 3.2

Under condition (C4), this follows by suitable extension of traditional arguments given for proving asymptotic normality of maximum likelihood estimators in fixed parametric models; see, for example, Lehmann (1983, chap. 6). In essence,

$$\left(\begin{matrix} \sqrt{n}(\hat{\theta}_S - \theta_0) \\ \sqrt{n}(\hat{\gamma}_S - \gamma_{0,S}) \end{matrix}\right) \doteq J_S^{-1}\left(\begin{matrix} \sqrt{n}\,\overline{U}_n \\ \sqrt{n}\,\overline{V}_{n,S} \end{matrix}\right)$$

$$\xrightarrow{d} \left(\begin{matrix} J^{00,S} & J^{01,S} \\ J^{10,S} & J^{11,S} \end{matrix}\right)\left(\begin{matrix} J_{01}\delta + M \\ \pi_S J_{11}\delta + N_S \end{matrix}\right),$$

with N_S denoting the vector of N_j's with $j \in S$.

We mention that Lemma 3.2 often may be proved to hold under weaker conditions than (C4) in cases where the log-density is concave in the parameters. This is important for the somewhat more difficult statements and proofs required when extending Lemmas 3.1–3.3 to regression models. Space does not allow to give the details here, but transparent proofs, under minimal conditions of the type $n^{-1/2}\max_{i \leq n}\|x_i\| \to 0$, may be given for log-concave models using the convexity arguments of Hjort and Pollard (1994).

Proof of Lemma 3.3

Using a delta method Taylor expansion for $\mu(\hat{\theta}_S, \hat{\gamma}_S) - \mu(\theta_0, \gamma_0 + \delta/\sqrt{n})$, in conjunction with Lemma 3.2, we easily establish that there is a limit distribution, which can be represented as

$$\Lambda_S = \left(\frac{\partial\mu}{\partial\theta}\right)^{\mathrm{t}} C_S + \left(\frac{\partial\mu}{\partial\gamma_S}\right)^{\mathrm{t}} D_S - \left(\frac{\partial\mu}{\partial\gamma}\right)^{\mathrm{t}}\delta.$$

(For the delta method, see, e.g., Barndorff-Nielsen and Cox 1989, chap. 2.) It is furthermore clear that Λ_S in this form is normal, and it is not difficult to work out valid expressions for the mean and variance and, hence, the limiting mean squared error. We will take the trouble to first derive certain simplified expressions for Λ_S, however, because these will be fruitful also for other purposes.

Using Lemma 3.2 in connection with expressions for the blocks of J^{-1}, one finds after some algebraic manipulations that

$$C_S = (J^{00,S}J_{01} + J^{01,S}\pi_S J_{11})\delta + J^{00,S}M + J^{01,S}N_S$$

$$= J_{00}^{-1}J_{01}(I - K^{1/2}H_S K^{-1/2})\delta + J_{00}^{-1}M$$

$$- J_{00}^{-1}J_{01}\pi_S^{\mathrm{t}} K_S \pi_S(N - J_{10}J_{00}^{-1}M)$$

$$= J_{00}^{-1}J_{01}(I - K^{1/2}H_S K^{-1/2})\delta + J_{00}^{-1}M$$

$$- J_{00}^{-1}J_{01}K^{1/2}H_S K^{-1/2}W,$$

whereas similarly

$$D_S = (J^{10,S}J_{01} + J^{11,S}\pi_S J_{11})\delta + J^{10,S}M + J^{11,S}N_S$$

$$= K_S \pi_S K^{-1}(\delta + W).$$

This leads to a fruitful expression for Λ_S in terms of a bias part and a zero-mean normal. Its mean is $b_S^{\mathrm{t}}\delta$, where

$$b_S = (I - K^{-1}\pi_S^{\mathrm{t}} K_S \pi_S)J_{10}J_{00}^{-1}\frac{\partial\mu}{\partial\theta} + (K^{-1}\pi_S^{\mathrm{t}} K_S \pi_S - I)\frac{\partial\mu}{\partial\gamma}$$

$$= (I - K^{-1/2}H_S K^{1/2})\omega.$$

Working similarly with the random part, we find the expression for Λ_S given in Lemma 3.3. Using the independence between M and W, which was noted before stating this lemma, we easily obtain the variance formula stated in the lemma.

Proof of Theorem 4.1

There is simultaneous convergence in distribution of all the $\sqrt{n}(\hat{\mu}_S - \mu_{\text{true}})$ jointly with D_n to the corresponding collection of Λ_S and D. This follows via arguments used to prove Lemma 3.3 and the fact that all limit variables can be expressed in terms of $(M^{\text{t}}, N^{\text{t}})^{\text{t}}$. Thus, there is also joint convergence in distribution of all $\sqrt{n}(\hat{\mu}_S - \mu_{\text{true}})$ with $c(S \mid D_n)$ to corresponding Λ_S with $c(S \mid D)$, in that $c(S \mid d)$ is almost continuous in d. Consequently,

$$\sqrt{n}(\hat{\mu} - \mu_{\text{true}}) = \sum_S c(S \mid D_n)\sqrt{n}(\hat{\mu}_S - \mu_{\text{true}}) \xrightarrow{d} \sum_S c(S \mid D)\Lambda_S.$$

The second expression for Λ follows with some effort, using that $c(S \mid D)$ sum to 1 for fixed D, in conjunction with the representation featured in Lemma 3.3. As mentioned earlier, M and $W = D - \delta$ turn out to be stochastically independent. Hence, Λ, given $D = d$, is a normal distribution, with

$$\text{Var}(\Lambda \mid d) = \left(\frac{\partial \mu}{\partial \theta}\right)^{\text{t}} J_{00}^{-1} \frac{\partial \mu}{\partial \theta} = \tau_0^2,$$

the minimal possible limit distribution variance for the estimators under consideration, and $E(\Lambda \mid d) = \omega^{\text{t}}\{\delta - G(d)^{\text{t}}d\}$. It follows that the limiting mean squared error of an arbitrary estimator in the class under study can be expressed as $E\Lambda^2 = \tau_0^2 + E[\omega^{\text{t}}\{\delta - G(D)^{\text{t}}D\}]^2$.

Proof of Proposition 9.1

We choose to work with the case of the full model, that is, $S = \{1, \ldots, q\}$, where we also are content to write $\hat{\theta}$ and $\hat{\delta}$ for $\hat{\theta}_{\text{full}}$ and $\hat{\delta}_{\text{full}}$ and so on. The general case can be handled quite similarly. Introduce

$$Q_n(s, t) = \frac{L_n(\hat{\theta} + s/\sqrt{n}, \gamma_0 + (\hat{\delta} + t)/\sqrt{n})}{L_n(\hat{\theta}, \gamma_0 + \hat{\delta}/\sqrt{n})}$$

$$= \frac{L_n(\hat{\theta} + s/\sqrt{n}, \hat{\gamma} + t/\sqrt{n})}{L_n(\hat{\theta}, \hat{\gamma})}.$$

Then, with Taylor expansion analysis, one sees that

$$\log Q_n(s, t) = -\frac{1}{2}\binom{s}{t}^{\text{t}} J_n \binom{s}{t} + O_p\left(n^{-1/2}\left\|\binom{s}{t}\right\|^3\right).$$

For a calculation needed in a moment, note that, for a symmetric positive-definite $(p + q) \times (p + q)$ matrix A,

$$\int \exp\left\{-\frac{1}{2}\binom{s}{t}^{\text{t}} A \binom{s}{t}\right\} ds$$

$$= (2\pi)^{p/2}|A|^{-1/2}|A^{11}|^{-1/2}\exp\left\{-\frac{1}{2}t^{\text{t}}(A^{11})^{-1}t\right\},$$

where A^{11} is the $q \times q$ lower right submatrix of A^{-1}; this follows from properties of the multinormal density. Substituting $\theta = \hat{\theta} + s/\sqrt{n}$ and $\delta = \hat{\delta} + t$ in the λ_n integral now leads to

$$\lambda_n = L_n(\hat{\theta}, \hat{\gamma})n^{-p/2}\int Q_n(s, t)\pi_0(\hat{\theta} + s/\sqrt{n})\pi(\hat{\delta} + t)\, ds\, dt$$

$$\doteq L_n(\hat{\theta}, \hat{\gamma})n^{-p/2}\pi_0(\hat{\theta})(2\pi)^{p/2}|J_n|^{-1/2}|J_n^{11}|^{-1/2}$$

$$\times \int \pi(\hat{\delta} + t)\exp\left\{-\tfrac{1}{2}t^{\text{t}}(J_n^{11})^{-1}t\right\} dt.$$

This proves the claims made.

In this proof we have glossed over certain technicalities that in a more careful proof need attention. These have to do with process convergence of $\log Q_n(s, t)$ over the space of functions of (s, t) defined over a compact region and with limiting the size and influence of $\log Q_n(s, t)$ outside such a compact region. We omit these details here, but refer to techniques and details provided in Hjort (1986), invented and developed there for a different but sufficiently similar problem.

Proof of Proposition 9.2

We start by reexpressing

$$\tilde{\mu}_S = \frac{\int \mu(\theta, \gamma_{0,S} + \delta_S/\sqrt{n})L_{n,S}(\theta, \gamma_{0,S} + \delta_S/\sqrt{n})\pi_0(\theta)\pi_S(\delta_S)\, d\theta\, d\delta_S}{\int L_{n,S}(\theta, \gamma_{0,S} + \delta_S/\sqrt{n})\pi_0(\theta)\pi_S(\delta_S)\, d\theta\, d\delta_S}$$

$$= \frac{\int \mu(\hat{\theta}_S + s/\sqrt{n}, \gamma_{0,S} + (\hat{\delta}_S + t)/\sqrt{n})Q_{n,S}(s, t)\pi_0(\hat{\theta}_S + s/\sqrt{n})\pi_S(+\hat{\delta}_S + t)\, ds\, dt}{\int Q_{n,S}(s, t)\pi_0(\hat{\theta}_S + s/\sqrt{n})\pi_S(\hat{\delta}_S + t)\, ds\, dt},$$

in terms of

$$Q_{n,S}(s, t) = \frac{L_{n,S}(\hat{\theta}_S + s/\sqrt{n}, \gamma_{0,S} + (\hat{\delta}_S + t)/\sqrt{n})}{L_{n,S}(\hat{\theta}_S, \hat{\gamma}_S)}$$

$$\doteq \exp\left\{-\frac{1}{2}\binom{s}{t}^{\text{t}} J_{n,S} \binom{s}{t}\right\},$$

in generalization of the Q_n process used in the proof of the previous proposition. Taylor expanding the μ term here wrt the first parameter gives $\mu(\hat{\theta}_S, \gamma_{0,S} + (\hat{\delta}_S + t)/\sqrt{n})$ plus $(\partial\mu/\partial\theta)(\hat{\theta}_S, \gamma_{0,S} + (\hat{\delta}_S + t)/\sqrt{n})$ times $s/\sqrt{n}$, and then integrating over s, as in Proposition 7.1, shows indeed that $\sqrt{n}(\tilde{\mu}_S - \bar{\mu}_S) \to_p 0$. A fact used here is that the integral of s times the limit of $Q_{n,S}(s, t)$, over s, is 0.

For the rest of the proof, we use Taylor expansion w.r.t. the second parameter, and find

$$\bar{\mu}_S = E\left\{\hat{\mu}_S + \left(\frac{\partial\mu}{\partial\gamma_S}\right)^{\text{t}}(\delta_S - \hat{\delta}_S)/\sqrt{n} \,\middle|\, \hat{\delta}_S\right\} + o_p(n^{-1/2}),$$

which when compared to Lemma 3.3 gives the required result.

[Received November 2002. Revised June 2003.]

REFERENCES

Barndorff-Nielsen, O. E., and Cox, D. R. (1989), *Asymptotic Techniques for Use in Statistics*, London: Chapman & Hall.

Berger, J. O. (1982), "Estimation in Continuous Exponential Families: Bayesian Estimation Subject to Risk Restrictions and Inadmissibility Results," in *Statistical Decision Theory and Related Topics III*, eds. J. O. Berger and S. S. Gupta, New York: Academic Press, pp. 109–141.

Bickel, P. J. (1981), "Minimax Estimation of the Mean of a Normal Distribution When the Parameter Space Is Restricted," *The Annals of Statistics*, 9, 1301–1309.

—— (1983), "Minimax Estimation of the Mean of a Normal Distribution Subject to Doing Well at a Point," *Recent Advances in Statistics, Festschrift for Herman Chernoff*, eds. M. H. Rizvi, J. Rustagi, and D. Siegmund, New York: Academic Press, pp. 511–528.

—— (1984), "Parametric Robustness: Small Biases Can Be Worthwhile," *The Annals of Statistics*, 12, 864–879.

Breiman, L. (2001), "Statistical Modeling: The Two Cultures" (with discussion), *Statistical Science*, 16, 199–231.

Buckland, S. T., Burnham, K. P., and Augustin, N. H. (1997), "Model Selection: An Integral Part of Inference," *Biometrics*, 53, 603–618.

Bühlmann, P. (1999), "Efficient and Adaptive Post-Model-Selection Estimators," *Journal of Statistical Planning and Inference*, 79, 1–9.

Burnham, K. P., and Anderson, D. R. (2002), *Model Selection and Multimodel Inference: A Practical Information-Theoretic Approach* (2nd ed.), New York: Springer-Verlag.

Chatfield, C. (1995), "Model Uncertainty, Data Mining and Statistical Inference," *Journal of the Royal Statistical Society*, Ser. A, 158, 419–466.

Claeskens, G., and Hjort, N. L. (2003), "The Focused Information Criterion," *Journal of the American Statistical Association*, 98, 900–916.

Draper, D. (1995), "Assessment and Propagation of Model Uncertainty" (with discussion), *Journal of the Royal Statistical Society*, Ser. B, 57, 45–97.

Foster, D. P., and George, E. I. (1994), "The Risk Inflation Criterion for Multiple Regression," *The Annals of Statistics*, 22, 1947–1975.

George, E. I. (1986a), "Minimax Multiple Shrinkage Estimation," *The Annals of Statistics*, 14, 188–205.

—— (1986b), "Combining Minimax Shrinkage Estimators," *Journal of the American Statistical Association*, 81, 437–445.

Green, P. J. (2003), "Trans-Dimensional Markov Chain Monte Carlo" (with discussion), in *Highly Structured Stochastic Systems*, eds. P. J. Green, N. L. Hjort, and S. Richardson, London: Oxford University Press, pp. 179–206.

Hjort, N. L. (1986), "Bayes Estimators and Asymptotic Efficiency in Parametric Counting Process Models," *Scandinavian Journal of Statistics*, 13, 63–85.

Hjort, N. L., and Pollard, D. (1994), "Asymptotics for Minimizers of Convex Processes," statistical research report, University of Oslo, Dept. of Mathematics.

Hoeting, J. A., Madigan, D., Raftery, A. E., and Volinsky, C. T. (1999), "Bayesian Model Averaging: A Tutorial" (with discussion), *Statistical Science*, 19, 382–417. [A version where the number of misprints has been significantly reduced is available at *http://www.stat.washington.edu/raftery/.*]

Hosmer, D. W., and Lemeshow, S. (1989), *Applied Logistic Regression*, New York: Wiley.

Hurvich, C. M., and Tsai, C.-L. (1989), "Regression and Time Series Model Selection in Small Samples," *Biometrika*, 76, 297–307.

———— (1990), "The Impact of Model Selection on Inference in Linear Regression," *The American Statistician*, 44, 214–217.

Kabaila, P. (1995), "The Effect of Model Selection on Confidence Regions and Prediction Regions," *Econometric Theory*, 11, 537–549.

———— (1998), "Valid Confidence Intervals in Regression After Variable Selection," *Econometric Theory*, 14, 463–482.

Leeb, H., and Pötscher, B. M. (2000), "The Finite-Sample Distribution of Post-Model-Selection Estimators, and Uniform Versus Non-Uniform Approxima-

tions," Technical Report TR 2000-03, Universität Wien, Institut für Statistik und Decision Support Systems.

Lehmann, E. L. (1983), *Theory of Point Estimation*, New York: Wiley.

Pötscher, B. M. (1991), "Effects of Model Selection on Inference," *Econometric Theory*, 7, 163–185.

Rao, J. S., and Tibshirani, R. (1997), "The Out-Of-Bootstrap Method for Model Averaging and Selection," technical report, University of Toronto, Dept. of Statistics.

Schwarz, G. (1978), "Estimating the Dimension of a Model," *The Annals of Statistics*, 6, 461–464.

Sen, P. K., and Saleh, A. K. M. E. (1987), "On Preliminary Test and Shrinkage *M*-Estimation in Linear Models," *The Annals of Statistics*, 15, 1580–1592.

Serfling, R. (1980), *Approximation Theorems of Mathematical Statistics*, New York: Wiley.

Spiegelhalter, D. J., Best, N. G., Carlin, B. P., and van der Linde, A. (2002), "Bayesian Measures of Model Complexity and Fit" (with discussion), *Journal of the Royal Statistical Society*, Ser. B, 64, 583–639.

Yang, Y. (2001), "Adaptive Regression by Mixing," *Journal of the American Statistical Association*, 96, 574–588.